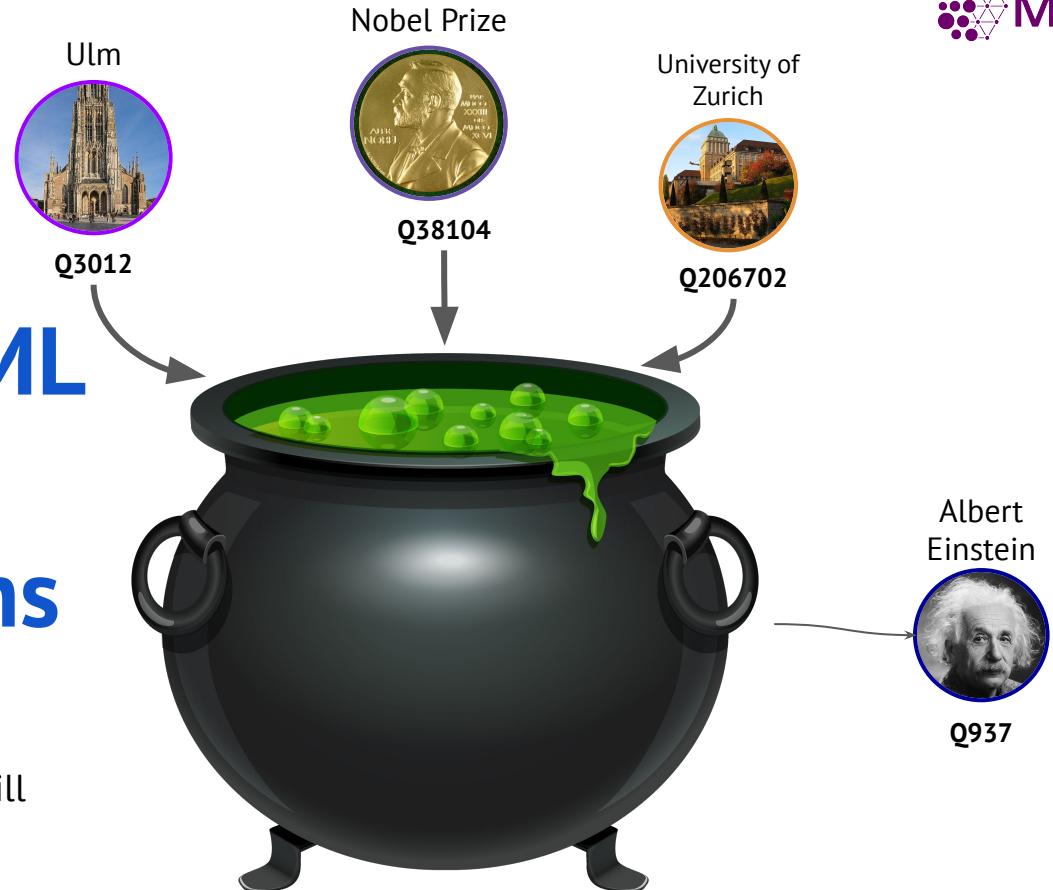


Discovering new ML applications on Knowledge Graphs

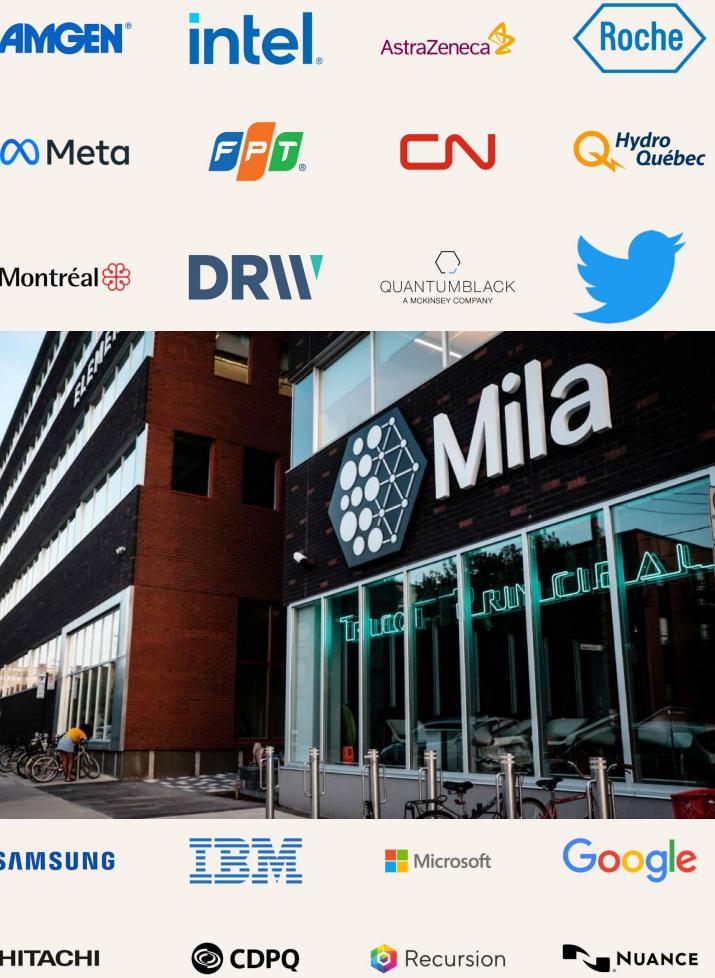
Michael Galkin
Postdoctoral Fellow @ Mila & McGill





About Mila

- Research Institute w/ UdeM, McGill, HEC Montreal, Polytechnique Montreal
 - Core Machine Learning, Deep Learning, Reinforcement Learning, NLP, Graph Learning
 - 100+ papers / year at top ML/AI conferences
- 900+ student researchers
- 40+ professors and CIFAR chairs
- 70+ industry partners
- 40+ Mila-affiliated startups

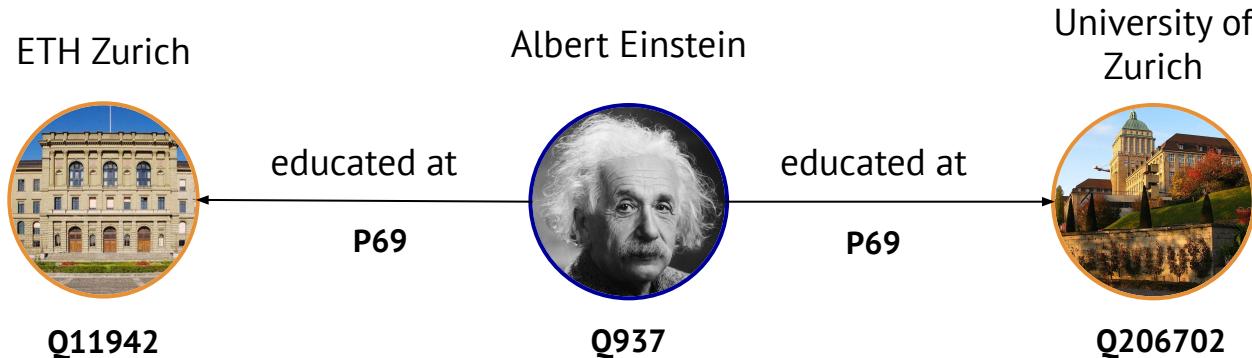


Plan

- **Vanilla KG Representation Learning: Re-cap**
- The New Big Picture
- NodePiece: Beyond Shallow Embeddings
- Hyper-Relational KGs
- Inductive Link Prediction with HR KGs
- Complex Query Answering with HR KGs
- Past, Today, Future



Triple-based Knowledge Graphs



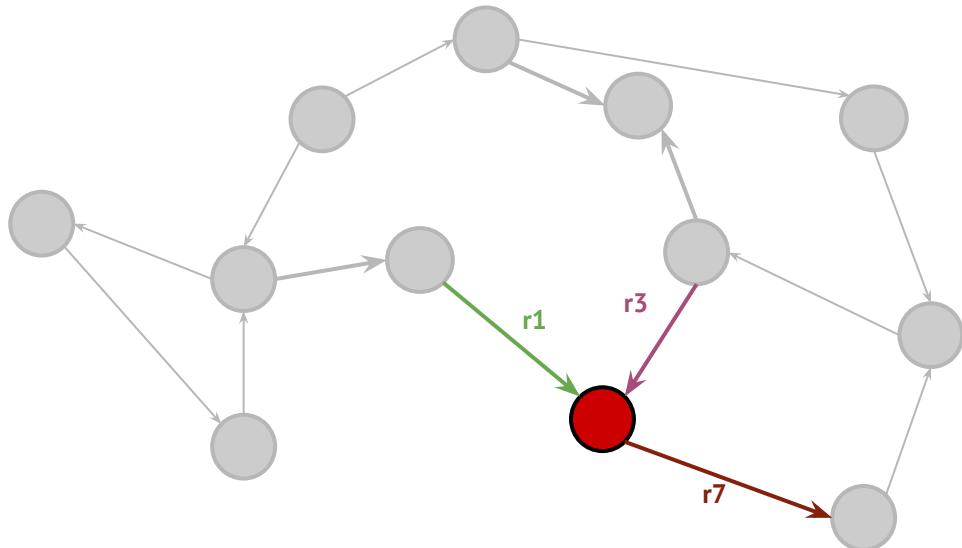
RDF

Albert Einstein
Albert Einstein

educatedAt
educatedAt

University of Zurich .
ETH Zurich .

Knowledge Graphs: Setup



- Directed graphs
- Explicit relation types (learnable edge features)
- Input node features are **not** given

Brief History: 2011 -

RESCAL

[Nickel et al, ICML 2011]

TransE

[Bordes et al, NeurIPS 2013]

100+ KG embedding models since then 😱



Brief History: 2011 -

Transductive

Triples

Supervised

RESCAL

[Nickel et al, ICML 2011]

TransE

[Bordes et al, NeurIPS 2013]

100+ KG embedding models since then 😱



Link Prediction on FB15k-237

Leaderboard Dataset



No substantial progress since 2018

Brief History: 2011 -

Transductive

Triples

Supervised

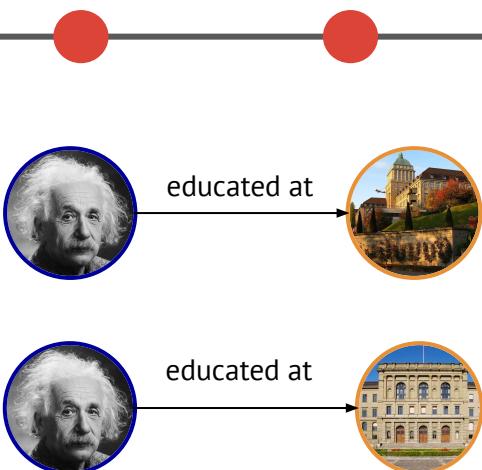
RESCAL

[Nickel et al, ICML 2011]

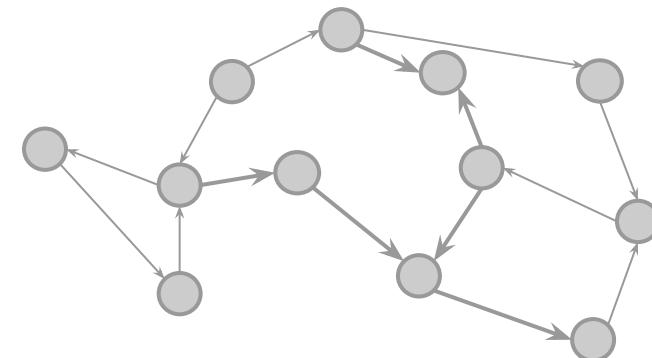
TransE

[Bordes et al, NeurIPS 2013]

100+ KG embedding models since then 😱



Can we go from
single triples to
encoding the whole
graph structure?



Brief History: 2011 -

Transductive

Triples

Supervised

RESCAL

[Nickel et al, ICML 2011]

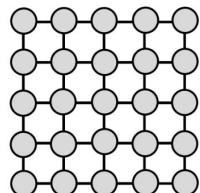
TransE

[Bordes et al, NeurIPS 2013]



The “5G” of Geometric Deep Learning

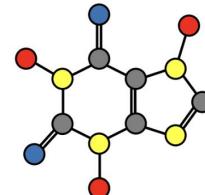
Geometric DL 
2018



Images &
Sequences



Homogeneous
spaces



Graphs & Sets



Manifolds, Meshes &
Geometric graphs

Geometric Deep Learning

Study of symmetries and invariances that unifies many deep learning architectures

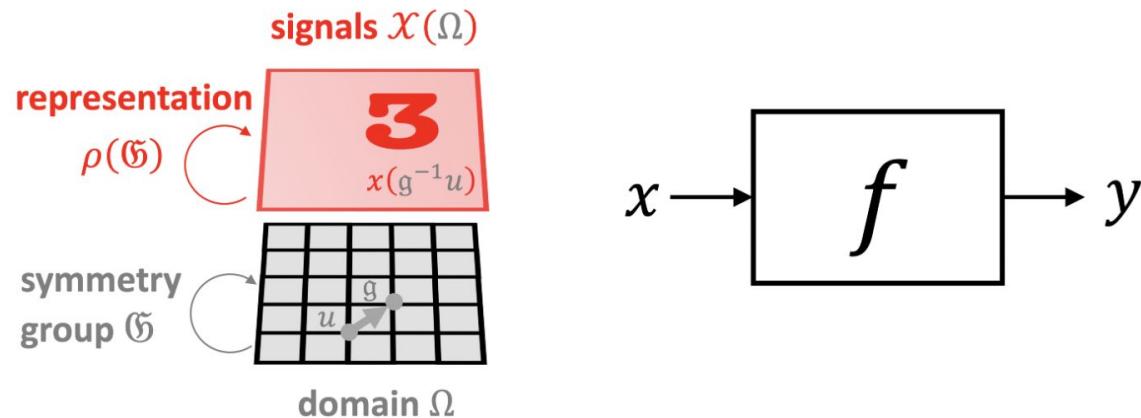
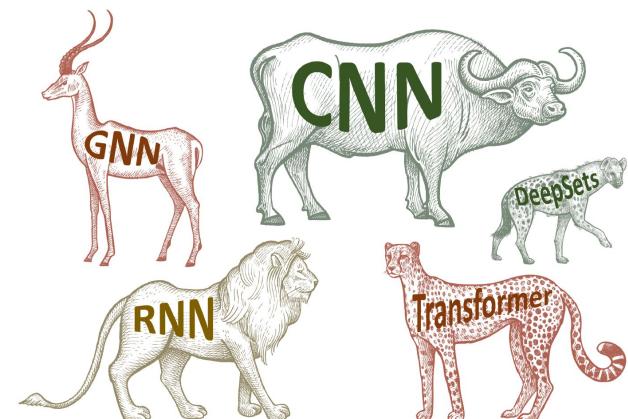


Illustration of geometric priors: the input signal (image $x \in \mathcal{X}(\Omega)$) is defined on the domain (grid Ω), whose symmetry (translation group \mathfrak{G}) acts in the signal space through the group representation $\rho(g)$ (shift operator).

Making an assumption on how the functions f (e.g. image classifier) interacts with the group restricts the hypothesis class.



Plan

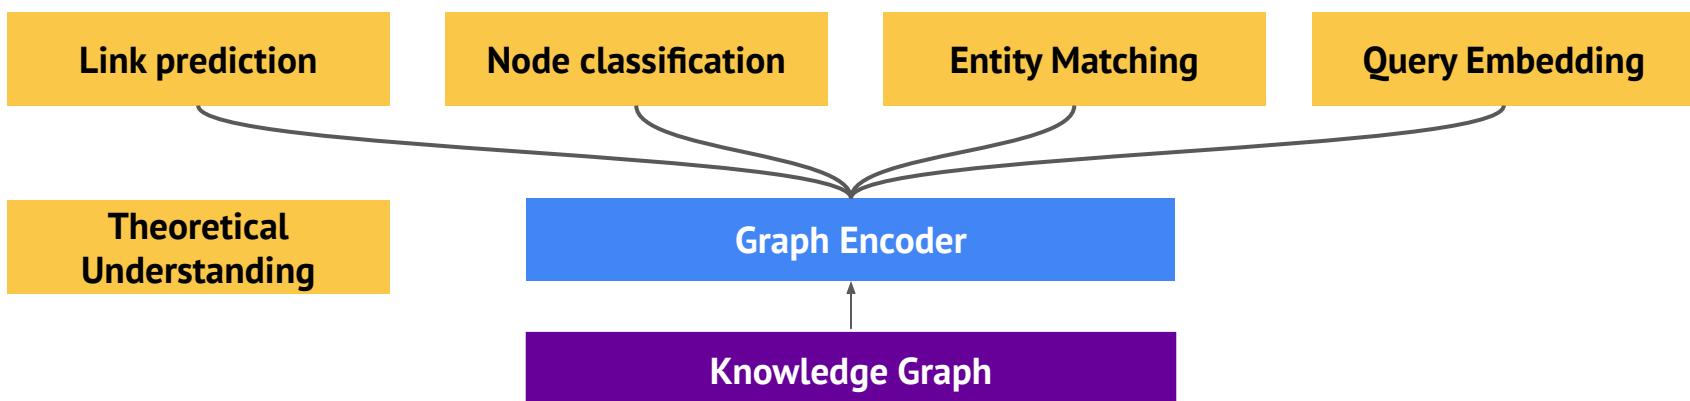
- Vanilla Triple KGs: Re-cap
- **The New Big Picture**
- NodePiece: Beyond Shallow Embeddings
- Hyper-Relational KGs
- Inductive Link Prediction with HR KGs
- Complex Query Answering with HR KGs
- Past, Today, Future

Big Picture in \mathbb{R}^5

Transductive	Triples	Supervised	Unimodal	Small
Inductive	Hyper-relational	Unsupervised	Multimodal	Large (sampling)

SETTING

TASK

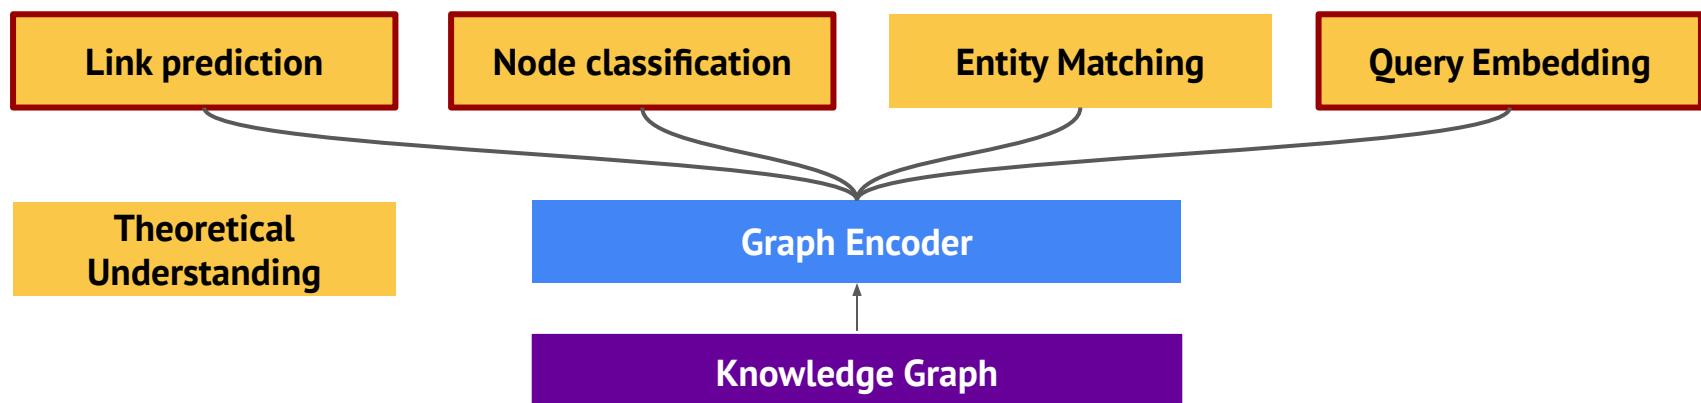


Big Picture in \mathbb{R}^5



SETTING

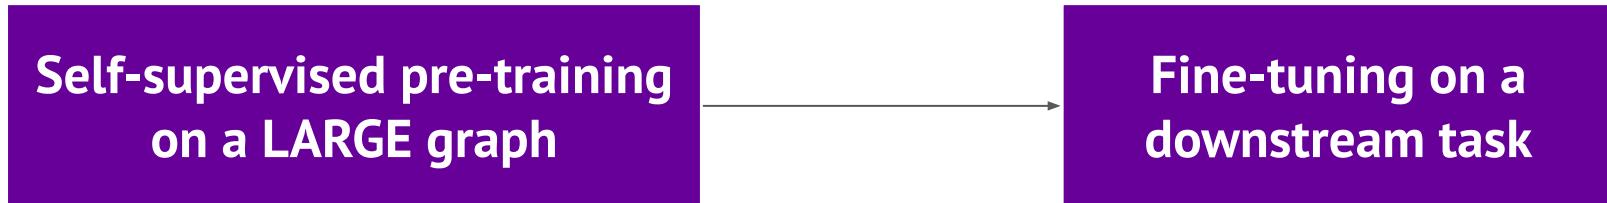
TASK



Plan

- Vanilla Triple KGs: Re-cap
- The New Big Picture
- **NodePiece: Beyond Shallow Embeddings**
- Hyper-Relational KGs
- Inductive Link Prediction with HR KGs
- Complex Query Answering with HR KGs
- Past, Today, Future

The ImageNet Moment for KGs



Graph ML

NLP

Vision

The ImageNet Moment for KGs

Self-supervised pre-training
on a LARGE graph

Fine-tuning on a
downstream task

Wikidata: 100M nodes

Embs: [100M, dim] ?

PyTorch BigGraph

~200 GB



Graph ML

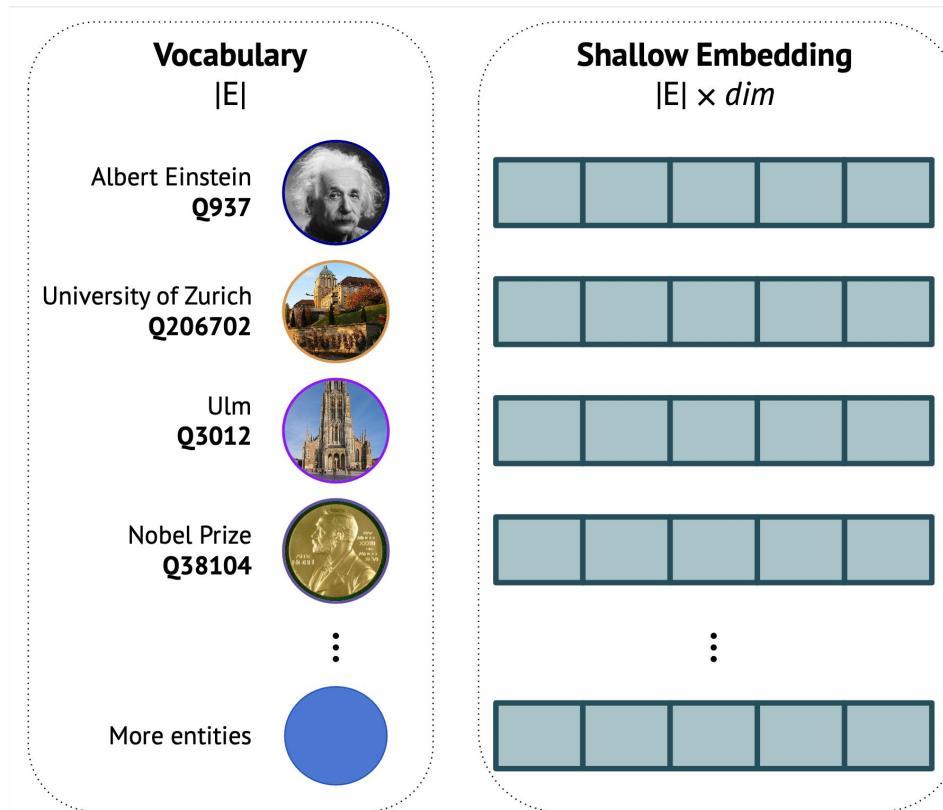
NLP

Vision

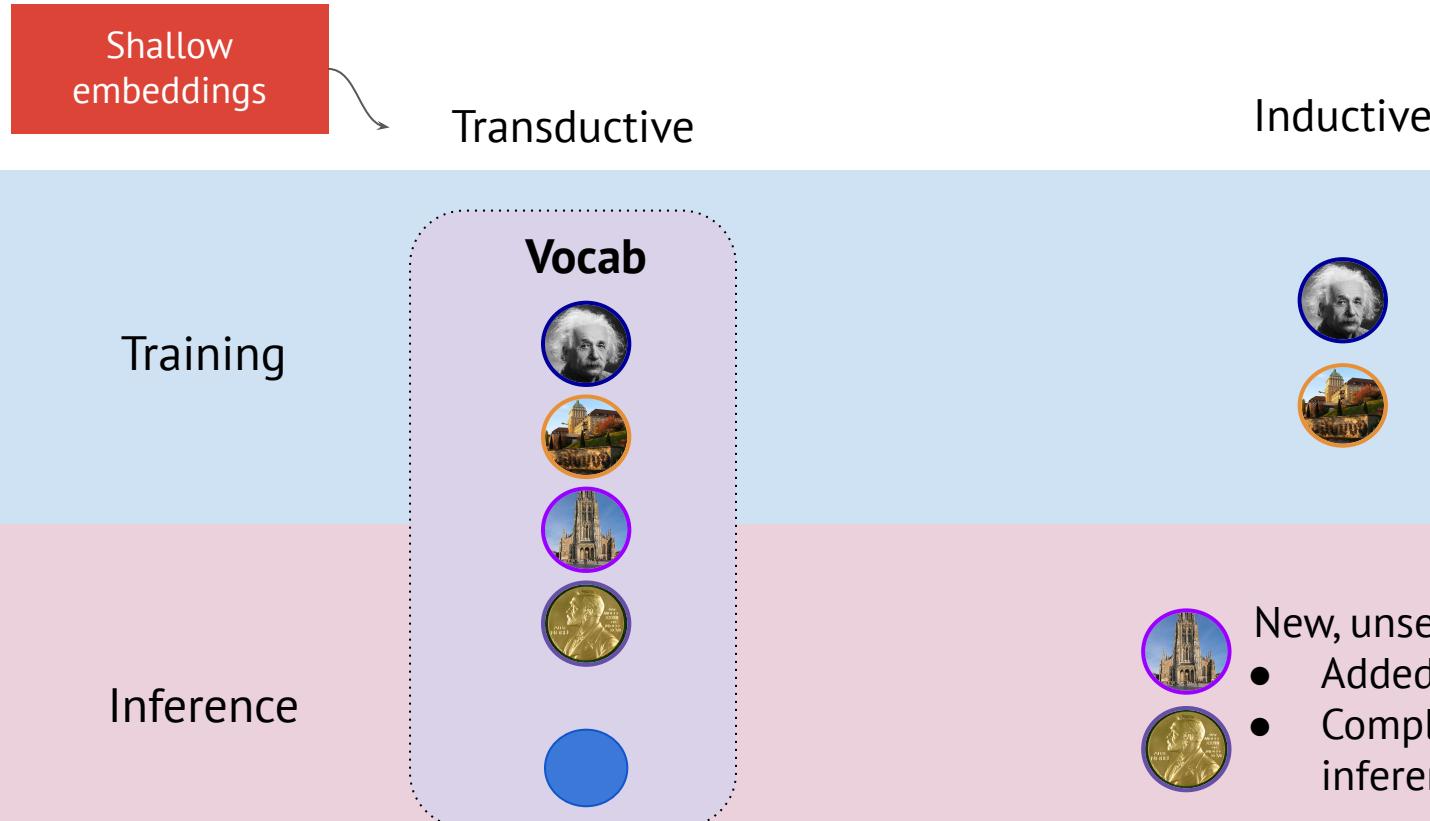
Shallow Embedding

Looks like a
Representation
Learning challenge 🤔

Can we do better?



Transductive vs Inductive



OGB WikiKG: Just 2.5M nodes (June'21)

Leaderboard for [ogbl-wikikg2](#)

The MRR score on the test and validation sets. The higher, the better.

Package: $\geq 1.2.4$

Deprecated [ogbl-wikikg](#) leaderboard can be found [here](#).

BERT-Large is $\sim 340M$ params

Rank	Method	Test MRR	Validation MRR	Contact	References	#Params	Hardware	Date
1	PairRE (200dim)	0.5208 ± 0.0027	0.5423 ± 0.0020	Linlin Chao	Paper , Code	500,334,800	Tesla P100 (16GB GPU)	Jan 28, 2021
2	RotatE (250dim)	0.4332 ± 0.0025	0.4353 ± 0.0028	Hongyu Ren – OGB team	Paper , Code	1,250,435,750	Quadro RTX 8000 (45GB GPU)	Jan 23, 2021
3	TransE (500dim)	0.4256 ± 0.0030	0.4272 ± 0.0030	Hongyu Ren – OGB team	Paper , Code	1,250,569,500	Quadro RTX 8000 (45GB GPU)	Jan 23, 2021
4	ComplEx (250dim)	0.4027 ± 0.0027	0.3759 ± 0.0016	Hongyu Ren – OGB team	Paper , Code	1,250,569,500	Quadro RTX 8000 (45GB GPU)	Jan 23, 2021

BERT (340M params) - disruption in NLP 
KG embs (>1B params) - 😬

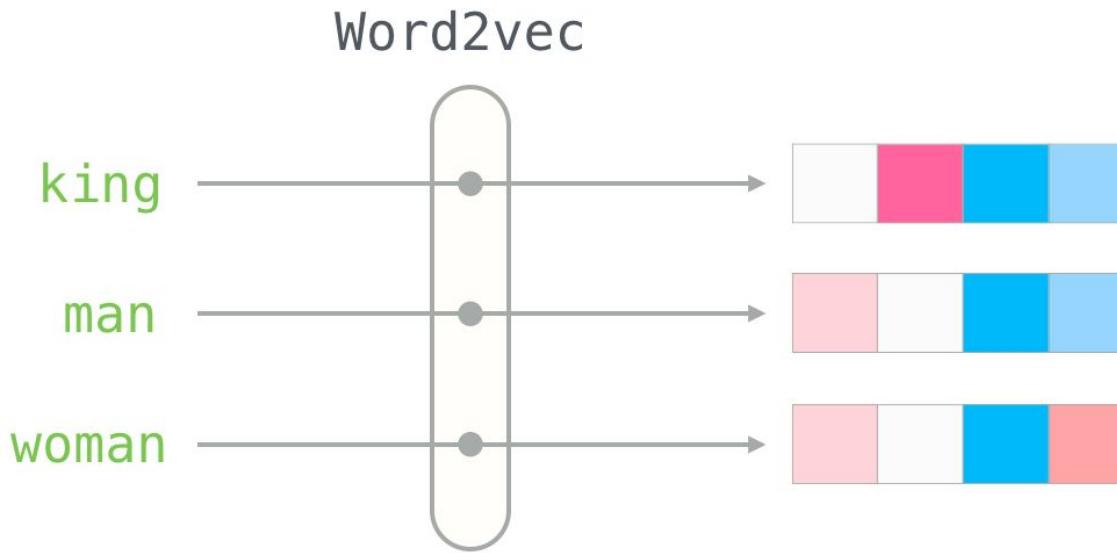
Life beyond shallow embedding?

Do we really need to learn & store the
whole **shallow** embedding matrix $|E| \times \text{dim}$?

Trying to fit a $100M \times 200$ tensor on a Tesla V100 ->



Back to 2014



Unseen words = [OOV] (out-of-vocabulary)

Byte-Pair Encoding / WordPiece

Dictionary

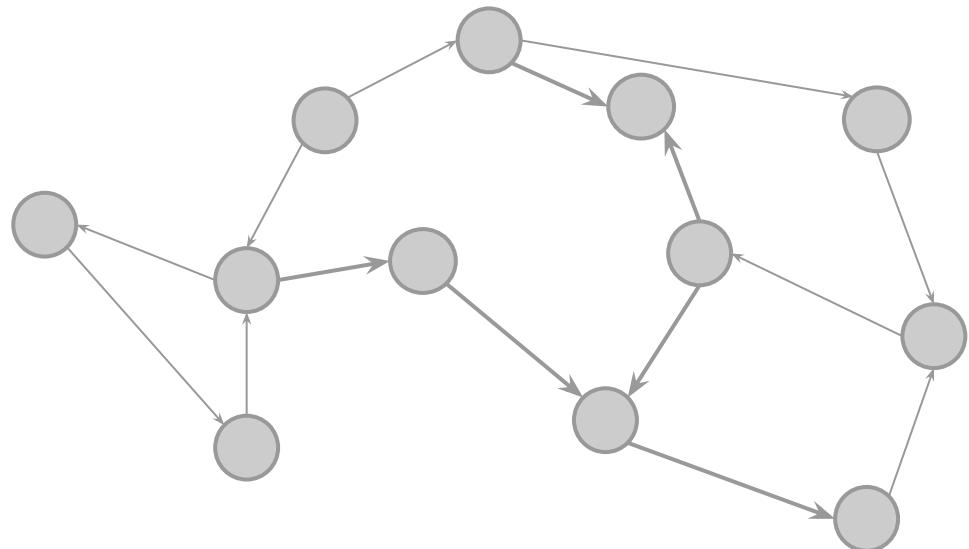
5 low
2 lower
6 newest
3 widest

Vocabulary

l, o, w, e, r, n, w, s, t, i, d, es, est

Add a pair (es, t) with freq 9

Tokenization + Graphs?

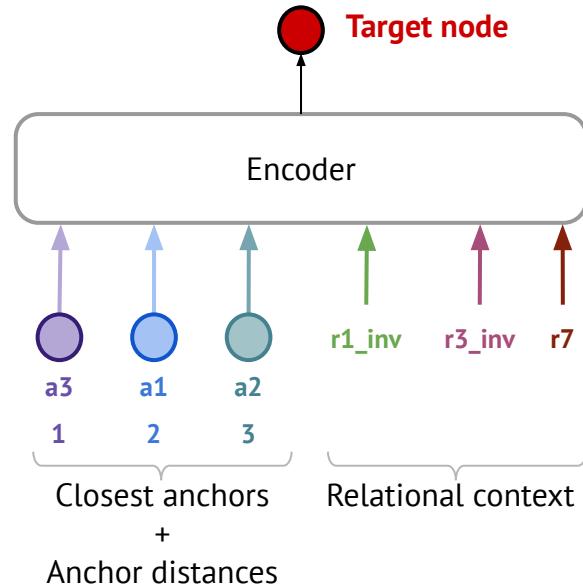
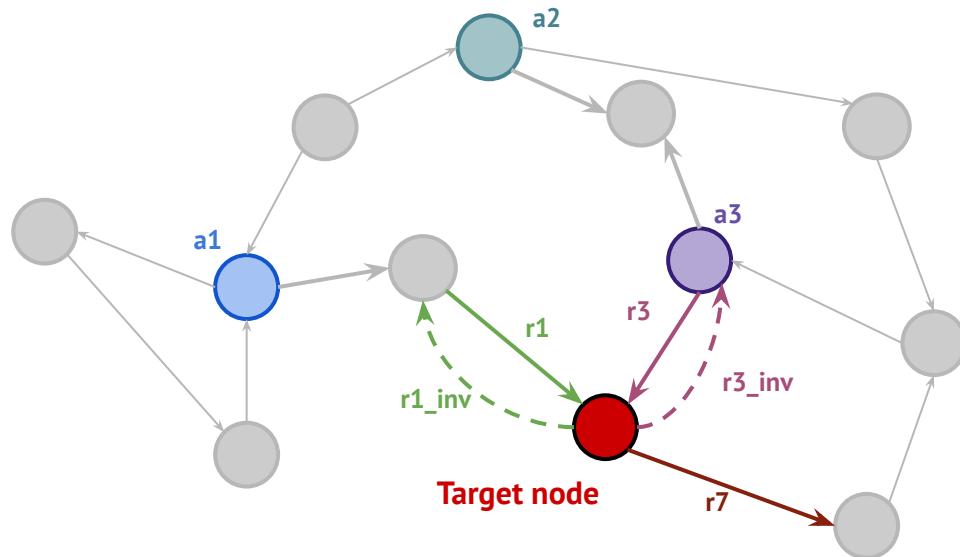


If nodes in a graph are
"words",
can we design a
fixed-size vocab of
"sub-word" units?

Tokenizing KGs

	Shallow embedding, only known words, otherwise OOV	Compositional representations, subword units
Language	Word2vec, GloVe	Byte-Pair Encoding, WordPiece
Graphs	All KG embedding algorithms (TransE, etc)	NodePiece

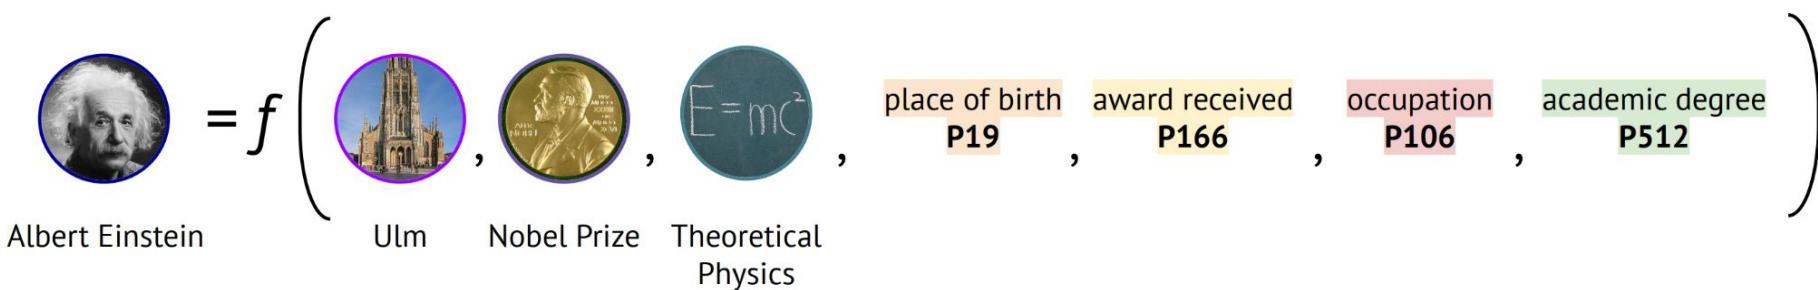
NodePiece - “*subword units*” for KGs



Vocabulary = Anchors + Relation types

Inductive out-of-the-box: unseen nodes are “tokenized” with the same Vocab

Tokenizing Einstein

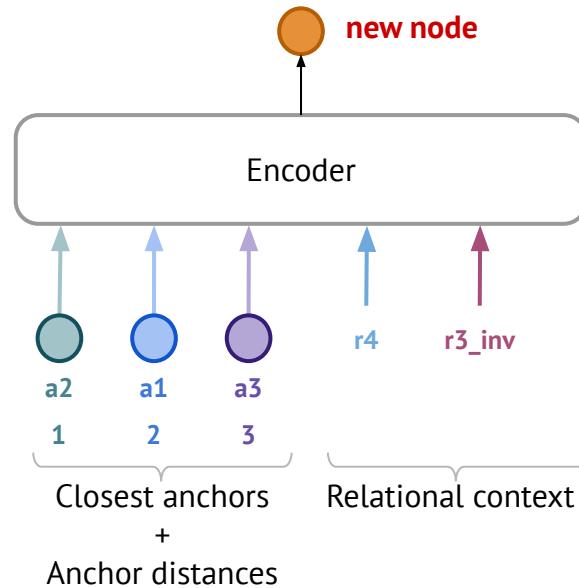
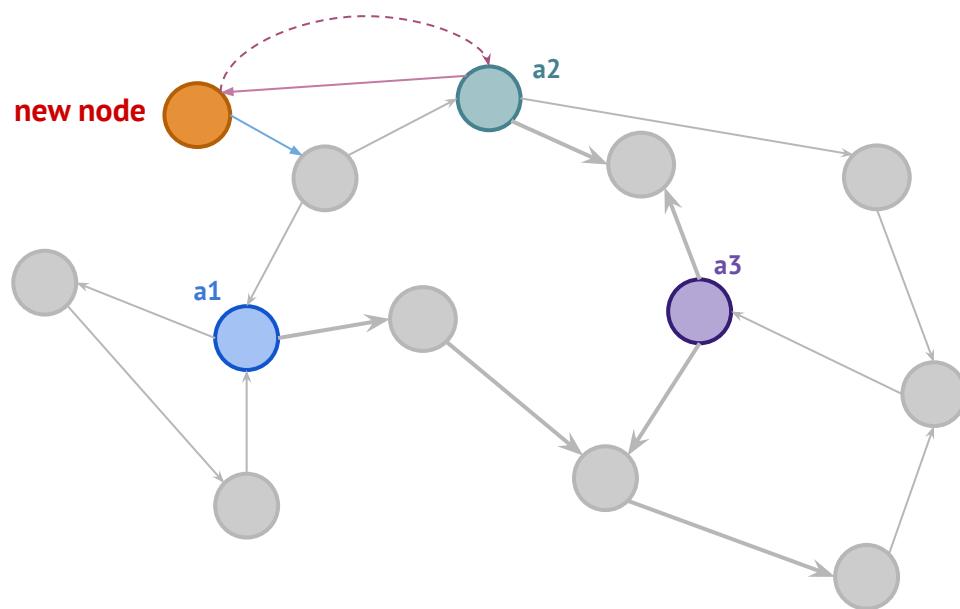


3 nearest anchors

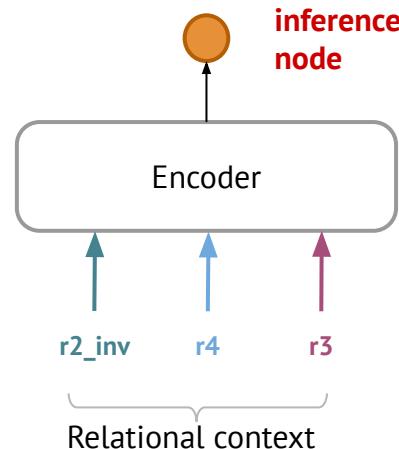
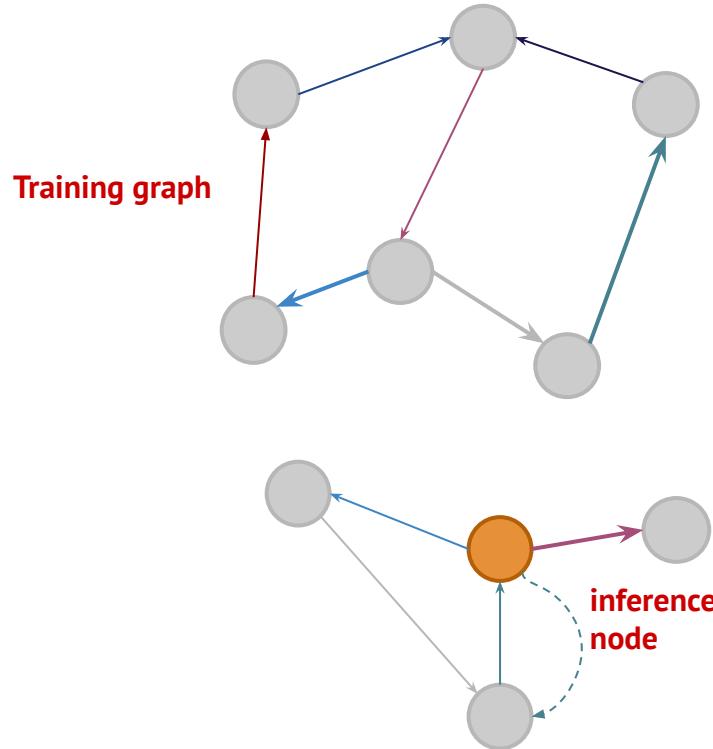
4 unique outgoing relations in the context



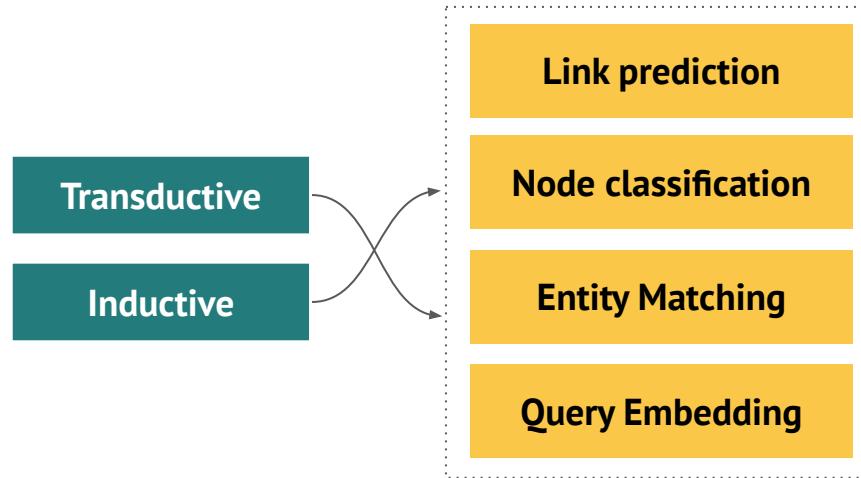
Unseen Node Tokenization



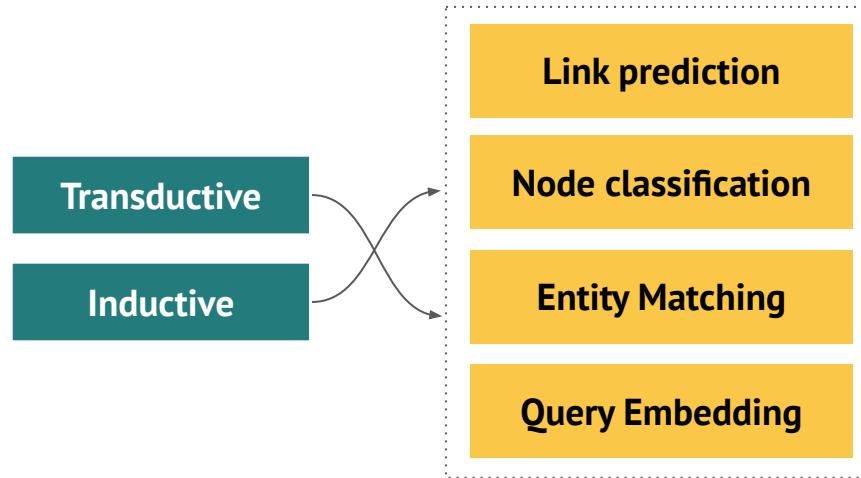
Inductive Node Tokenization



New Downstream Tasks



New Downstream Tasks



OGB WikiKG 2 : NodePiece is New SOTA

Leaderboard for [ogbl-wikikg2](#)

The MRR score on the test and validation sets. The higher, the better.

Package: $\geq 1.2.4$

Deprecated [ogbl-wikikg](#) leaderboard can be found [here](#).

February 2022

Rank	Method	Ext.		Validation			References	#Params	Hardware	Date
		data	Test MRR	MRR	Contact					
1	TripleRE + NodePiece	No	0.6866 ± 0.0014	0.6955 ± 0.0008	Long Yu (360AI)		Paper , Code	36,421,802	Tesla A100(40GB)	Feb 24, 2022
2	InterHT	No	0.6779 ± 0.0018	0.6893 ± 0.0015	Baoxin Wang (HFL)		Paper , Code	19,215,402	Tesla V100 (32GB)	Feb 10, 2022
3	TripleRE + NodePiece	No	0.6582 ± 0.0020	0.6616 ± 0.0018	Long Yu (360AI)		Paper , Code	7,289,002	Tesla A100(40GB)	Dec 25, 2021
4	ComplEx-RP (50dim)	No	0.6392 ± 0.0045	0.6561 ± 0.0070	Yihong Chen (UCL NLP & FAIR London)		Paper , Code	250,167,400	Tesla V100 (32GB)	Nov 23, 2021
5	tripleRE	No	0.5794 ± 0.0020	0.6045 ± 0.0024	Long Yu (360AI)		Paper , Code	500,763,337	Tesla P40(24GB)	Dec 17, 2021
6	NodePiece + AutoSF	No	0.5703 ± 0.0035	0.5806 ± 0.0047	Mikhail Galkin (Mila)		Paper , Code	6,860,602	Tesla V100 (32 GB)	Jul 17, 2021

NodePiece-enabled models

Inductive Link Prediction Challenge 2022

<https://github.com/pykeen/ilpc2022>

ILPC'22 Small

Split	Entities	Relations	Triples
Train	10,230	96	78,616
Inference	6,653	96 (subset)	20,960
Inference validation	6,653	96 (subset)	2,908
Inference test	6,653	96 (subset)	2,902
Hold-out test set	6,653	96 (subset)	2,894

ILPC'22 Large

Split	Entities	Relations	Triples
Train	46,626	130	202,446
Inference	29,246	130 (subset)	77,044
Inference validation	29,246	130 (subset)	10,179
Inference test	29,246	130 (subset)	10,184
Hold-out test set	29,246	130 (subset)	10,172

Model	MRR	H@100	H@10	H@5	H@3	H@1	AMRI
InductiveNodePieceGNN	0.1326	0.4705	0.2509	0.1899	0.1396	0.0763	0.730
InductiveNodePiece	0.0381	0.4678	0.0917	0.0500	0.0219	0.007	0.666

Plan

- Vanilla Triple KGs: Re-cap
- The New Big Picture
- NodePiece: Beyond Shallow Embeddings
- **Hyper-Relational KGs**
- **Inductive Link Prediction with HR KGs**
- **Complex Query Answering with HR KGs**
- Past, Today, Future

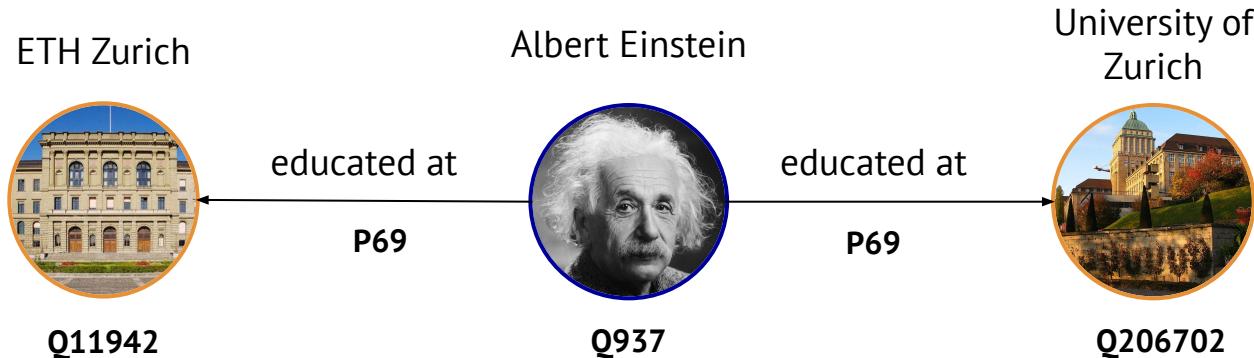


Triple-based Knowledge Graphs

Wait a sec

Studied the
same subject?

Got the same
degree?



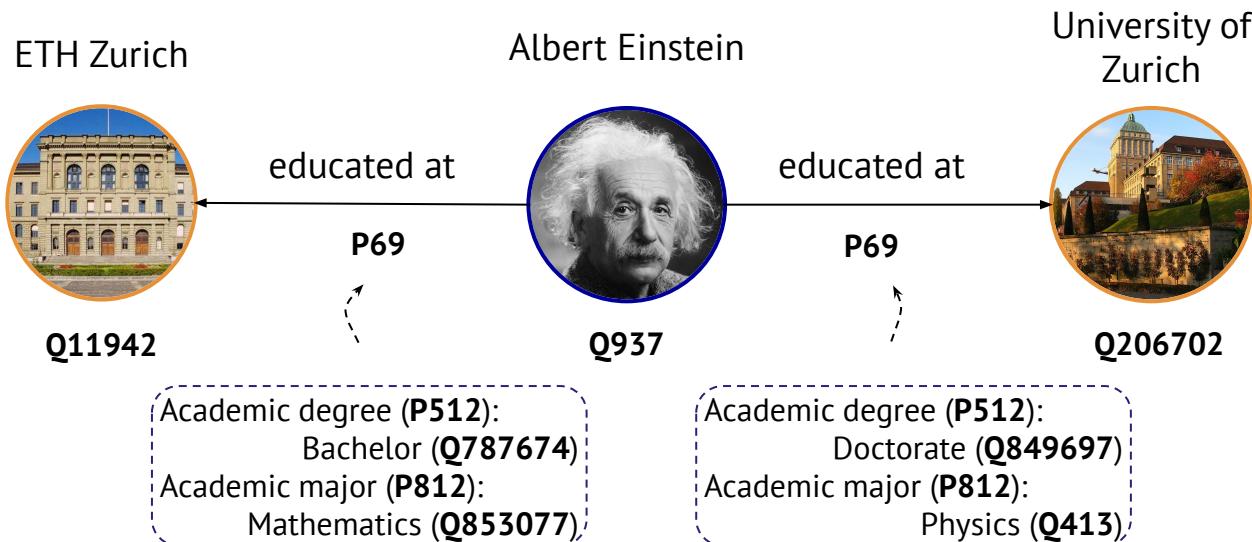
RDF

Albert Einstein
Albert Einstein

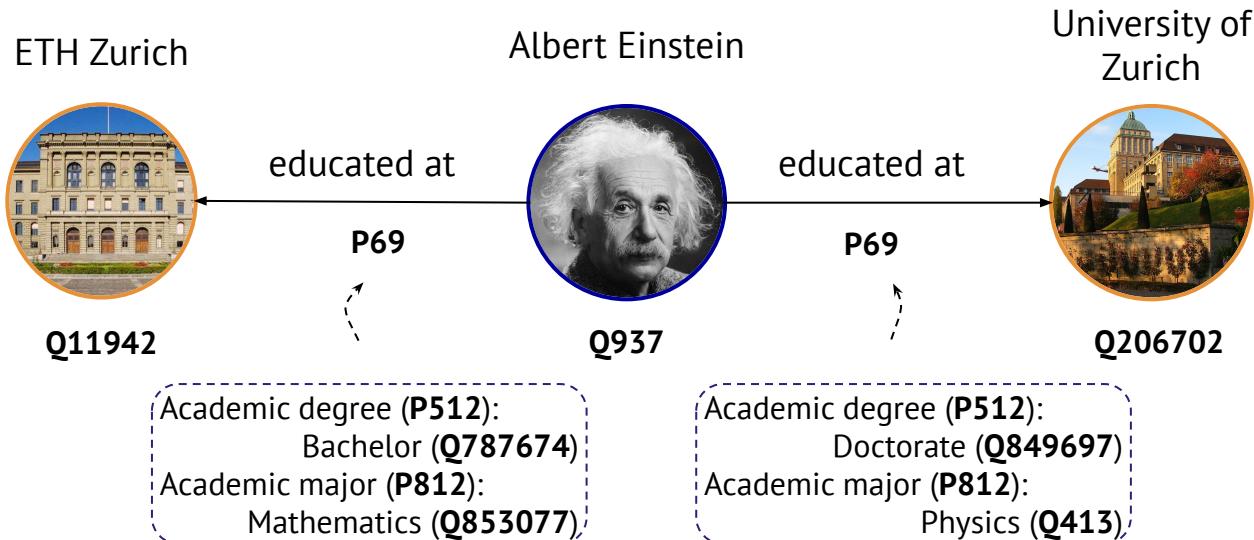
educatedAt
educatedAt

University of Zurich .
ETH Zurich .

Hyper-relational KGs



Hyper-relational KGs & RDF*



RDF*

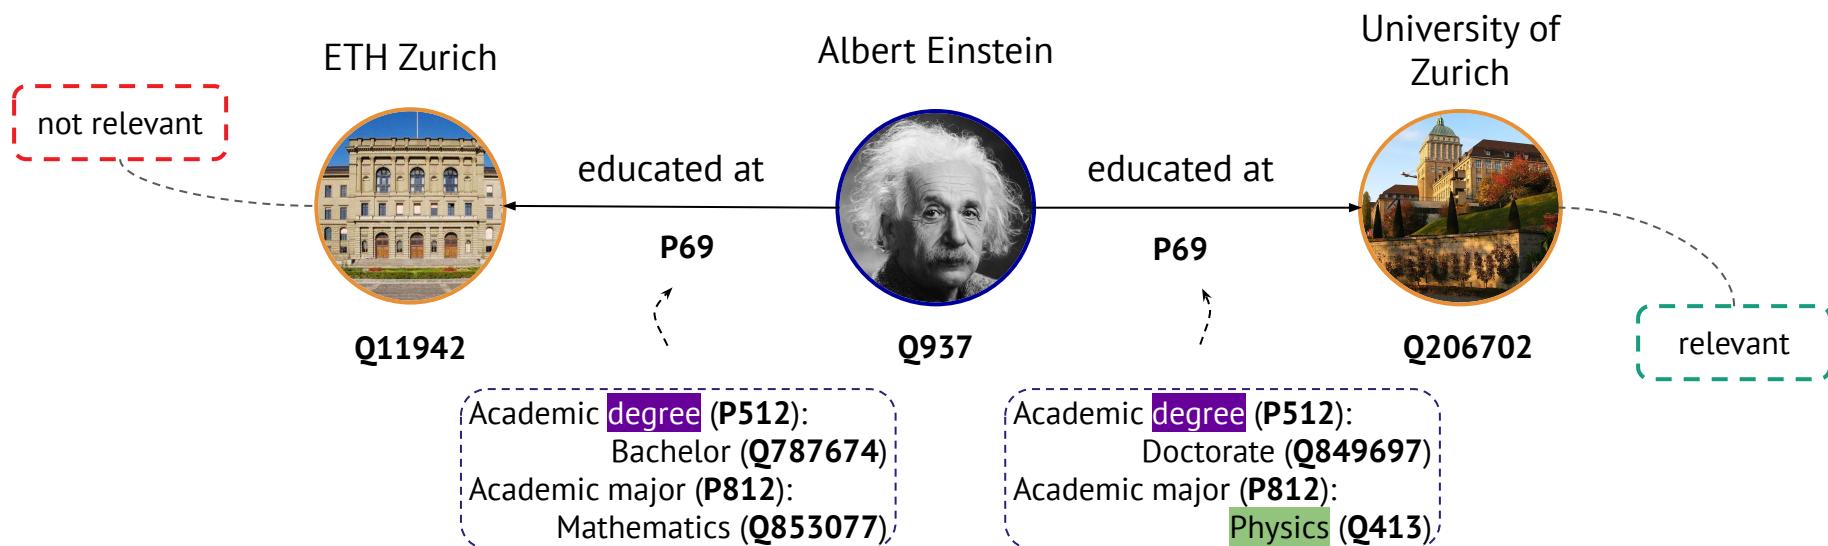
```

<< Albert Einstein educatedAt University of Zurich >>
    academic degree Doctorate;
    academic major Physics .
<< Albert Einstein educatedAt ETH Zurich >>
    Academic degree Bachelor ;
    Academic major Mathematics .
  
```

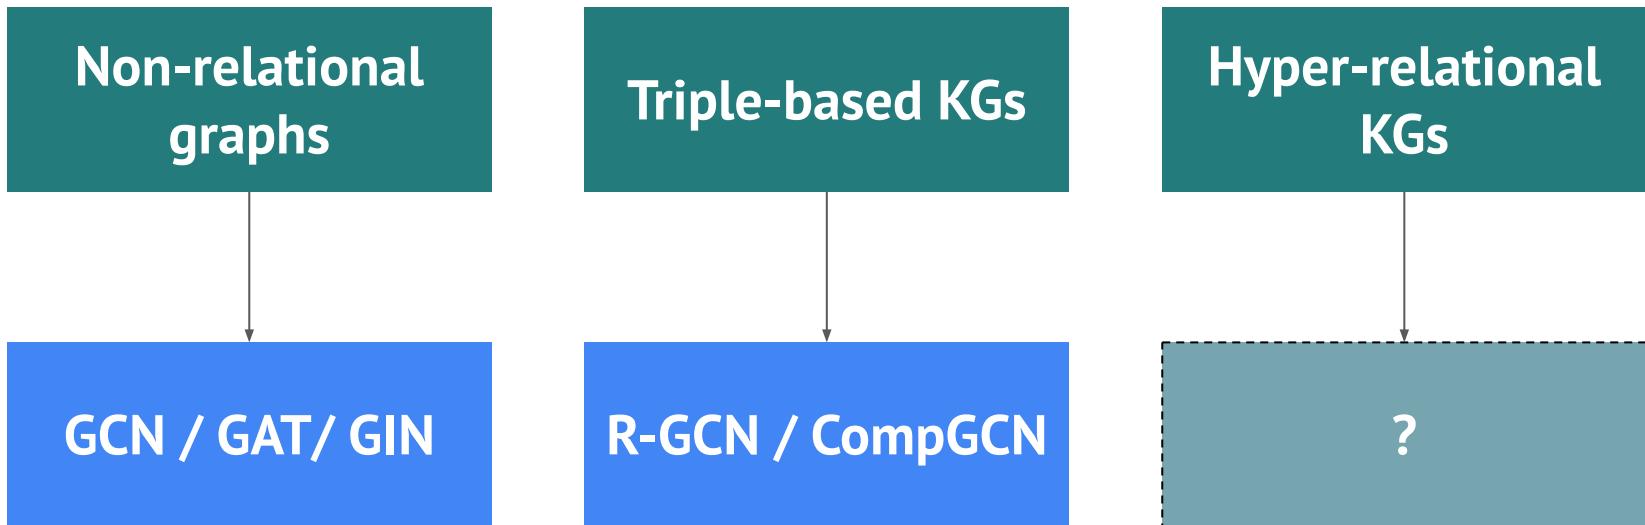
Statements with Qualifiers in Wikidata
Entity-relation Edge Attributes
Edge Instances

Hyper-relational KGs & RDF*

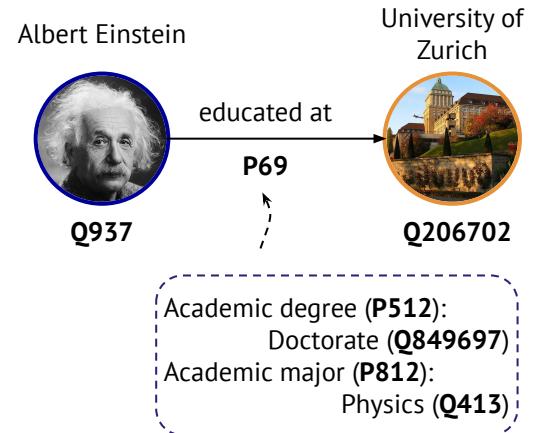
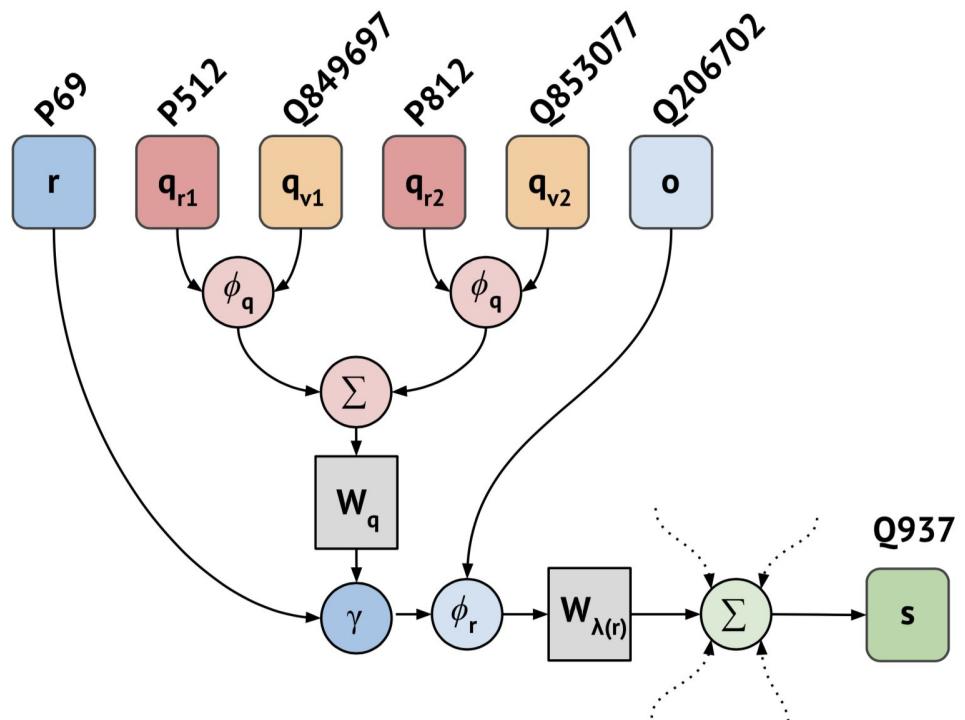
Where did Albert Einstein receive his degree in physics?



GNN Encoders

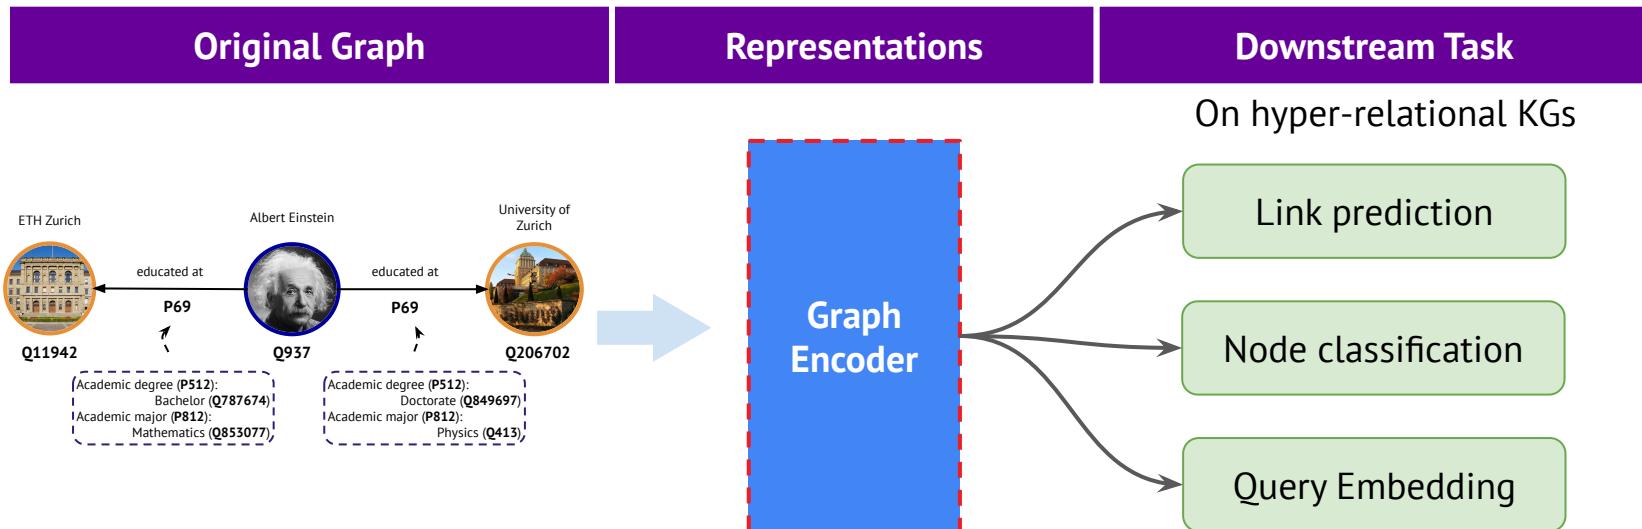


★ StarE: GNN architecture for HR KGs



$$\mathbf{h}_v = f \left(\sum_{(u,r) \in \mathcal{N}(v)} \mathbf{W}_{\lambda(r)} \phi_r(\mathbf{h}_u, \gamma(\mathbf{h}_r, \mathbf{h}_q)_{vu}) \right)$$

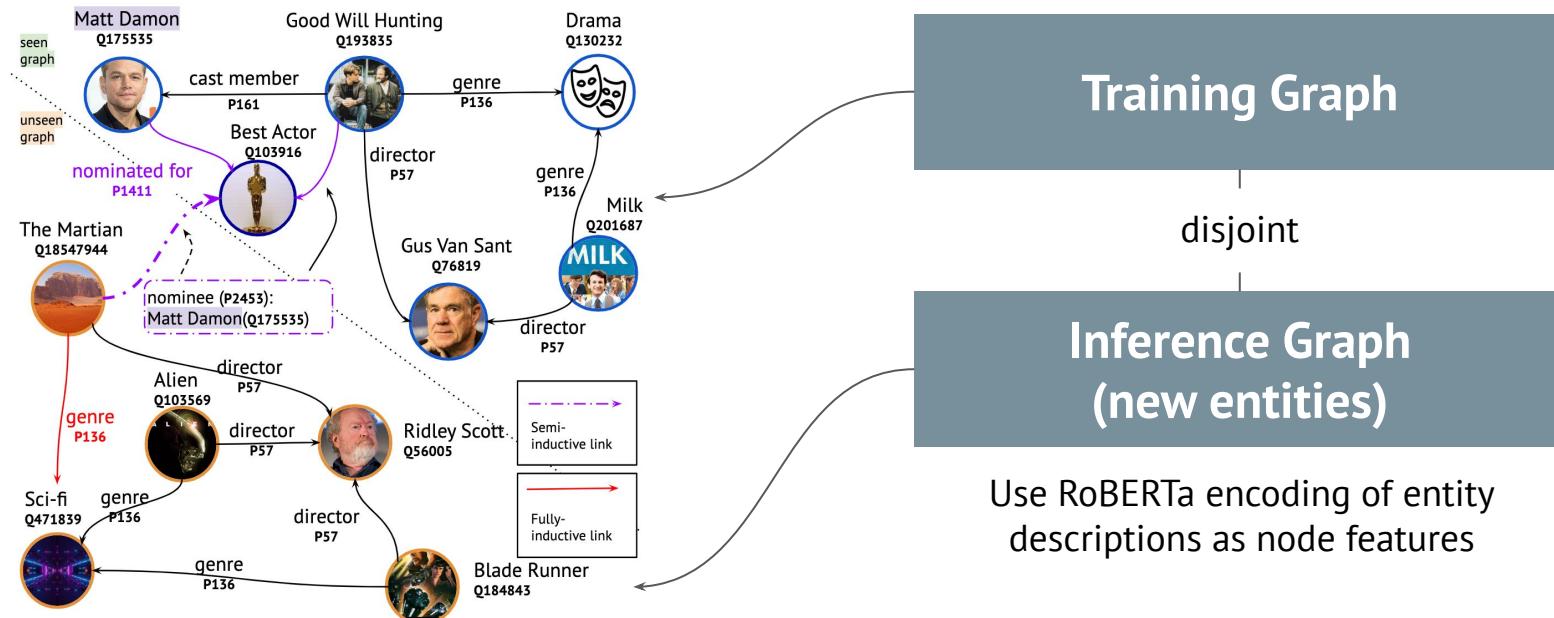
Enabling New Tasks



★ StarE Summary

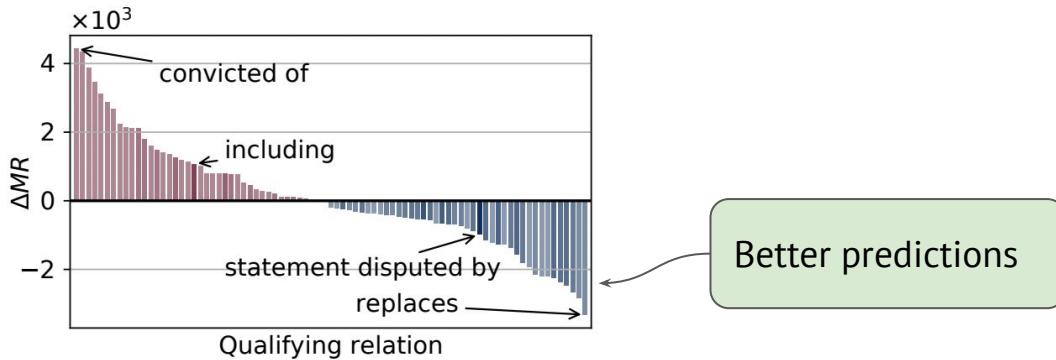
- A GNN encoder for embedding HR KGs and many downstream tasks
- Sparse Qualifier Representation
- As small as 1 qualifier per triple gives boosts
- The more qualifiers - the better

Hyper-relational KGs + Inductive Link Prediction



★ Inductive StarE Summary

- Qualifiers help in Inductive Link Prediction
- Features as encoded RoBERTa entity descriptions are good enough
- Some qualifiers give a lot of boost, some do not

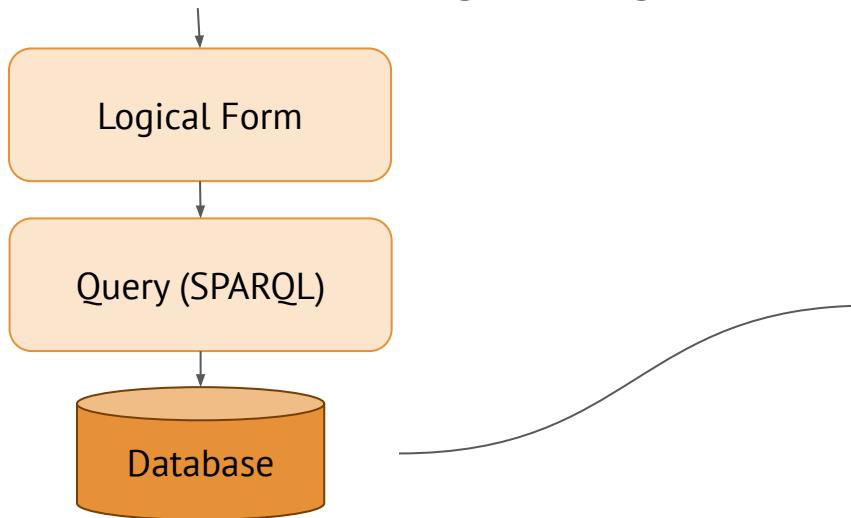


Complex Logical Query Answering: Why?



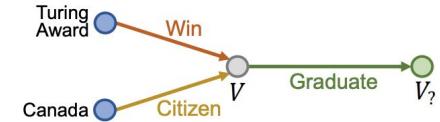
Where did Canadian citizens with Turing Award graduate?

Typical
KGQA
pipeline

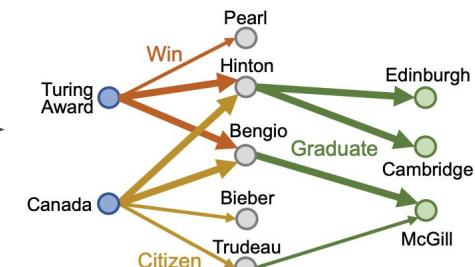


(A) Query q and Its Dependency Graph

$$q = V_? . \exists V : Win(TuringAward, V) \wedge Citizen(Canada, V) \wedge Graduate(V, V_?)$$



(C) Knowledge Graph Space

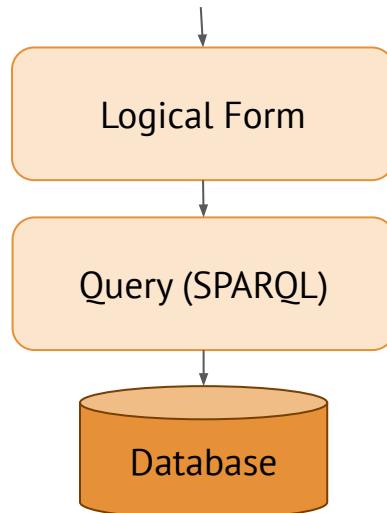


Complex Logical Query Answering: Why?



Where did Canadian citizens with Turing Award graduate?

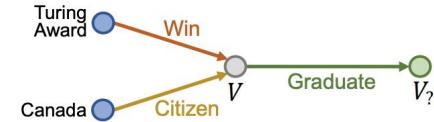
Typical KGQA pipeline



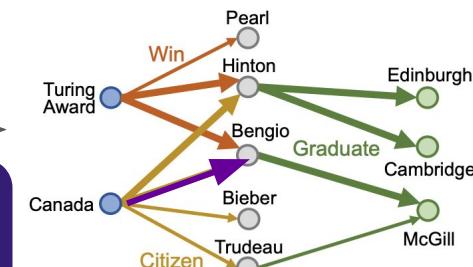
What if this edge is missing?

(A) Query q and Its Dependency Graph

$$q = V_? . \exists V : Win(TuringAward, V) \wedge Citizen(Canada, V) \wedge Graduate(V, V_?)$$



(C) Knowledge Graph Space



Complex Logical Query Answering

Transductive

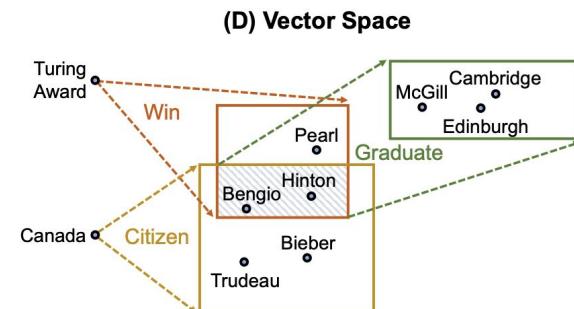
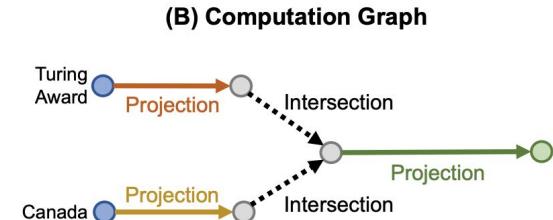
Triples

Inductive

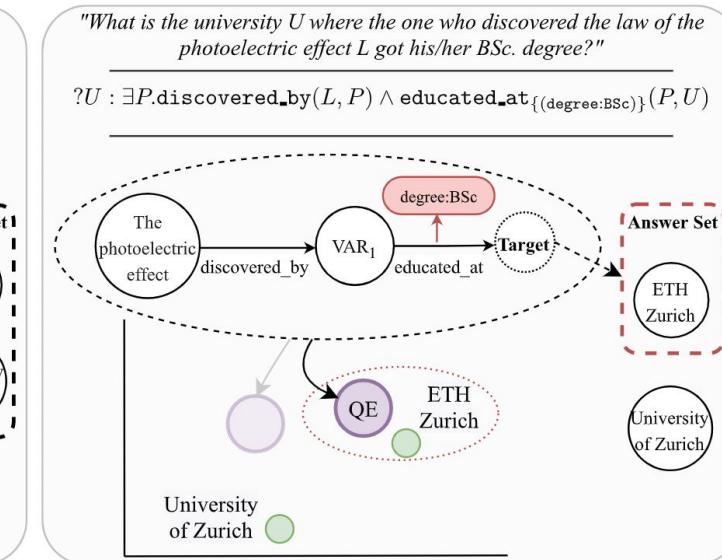
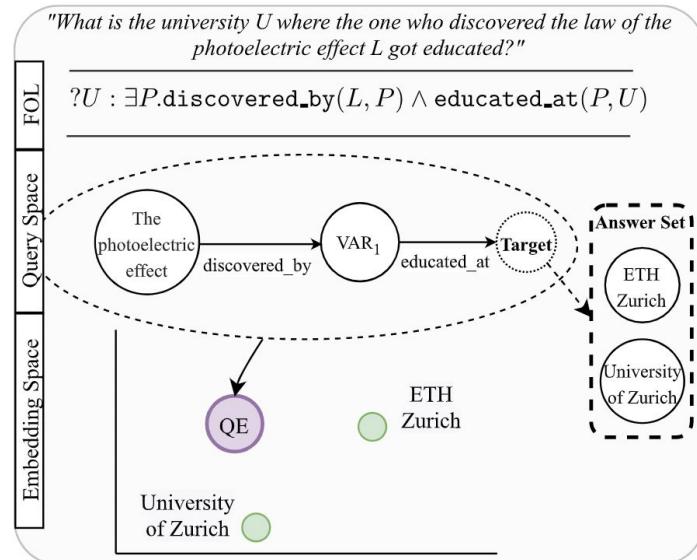
Hyper-relational

Query Embedding

- Databases assume KGs are **complete**
 - In reality - they are not
- We want to answer FOL queries over **incomplete** graphs with **neural** operators
- Embed a query in a latent space, MIPS decoder for kNN answers



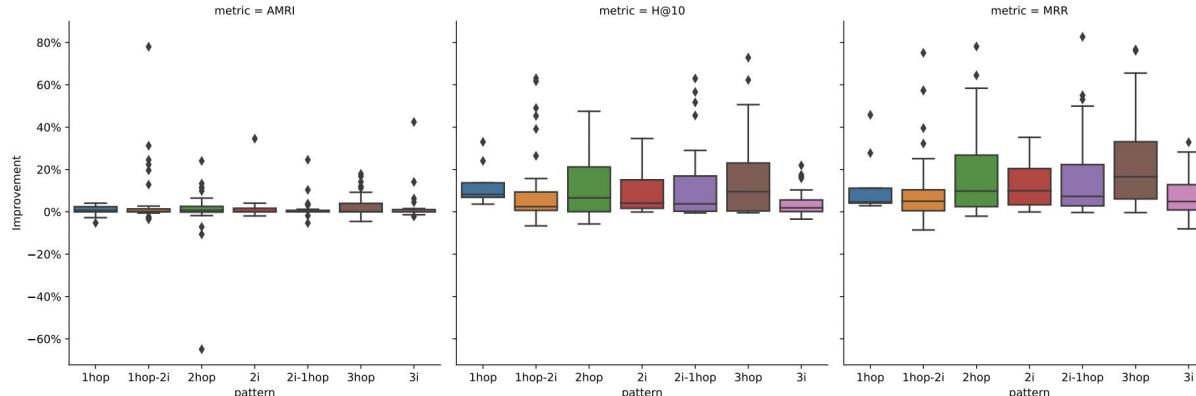
StarQE: Complex Logical Query Answering on HR KGs



1. Extending FOL to HR KGs
2. Qualifiers help A LOT
3. New query types are enabled

★ StarQE for Logical Queries: Summary

- Extend FOL to hyper-relational graphs with qualifiers
- Enabling new query types (eg, joins over qualifier entities)
- Robust to inner representation: RDF* vs reified RDF
- Qualifiers help A LOT in answering complex queries



Plan

- Vanilla Triple KGs: Re-cap
- The New Big Picture
- Hyper-Relational KGs
- Inductive Link Prediction with HR KGs
- Complex Query Answering with HR KGs
- **Past, Today, Future**

Space of KG Tasks in 2019

Transductive

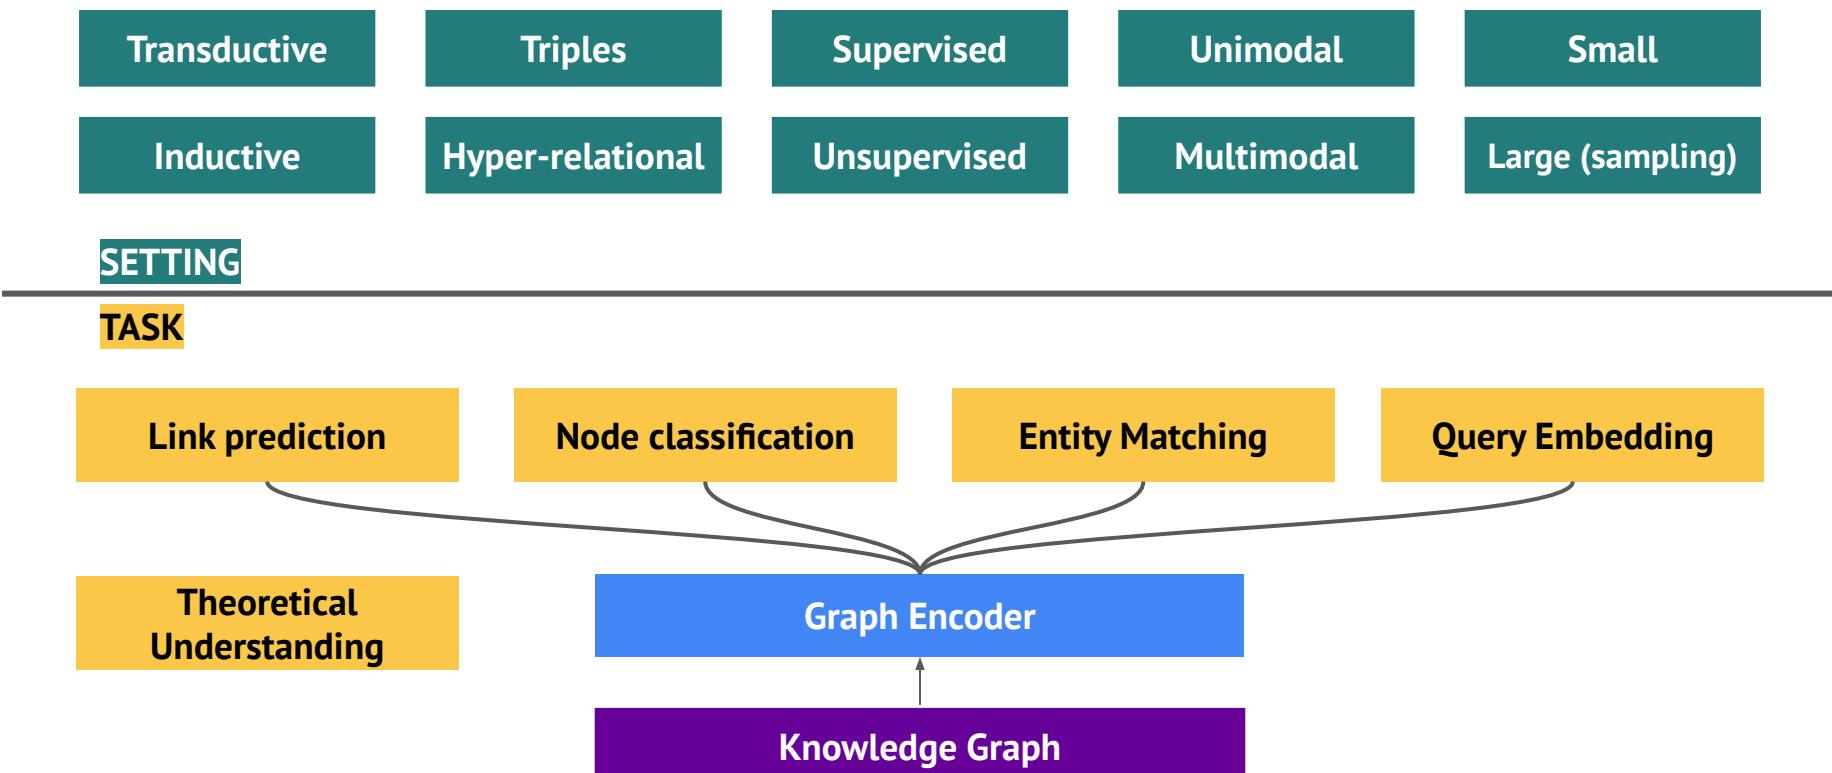
Triples

SETTING

TASK

Link prediction

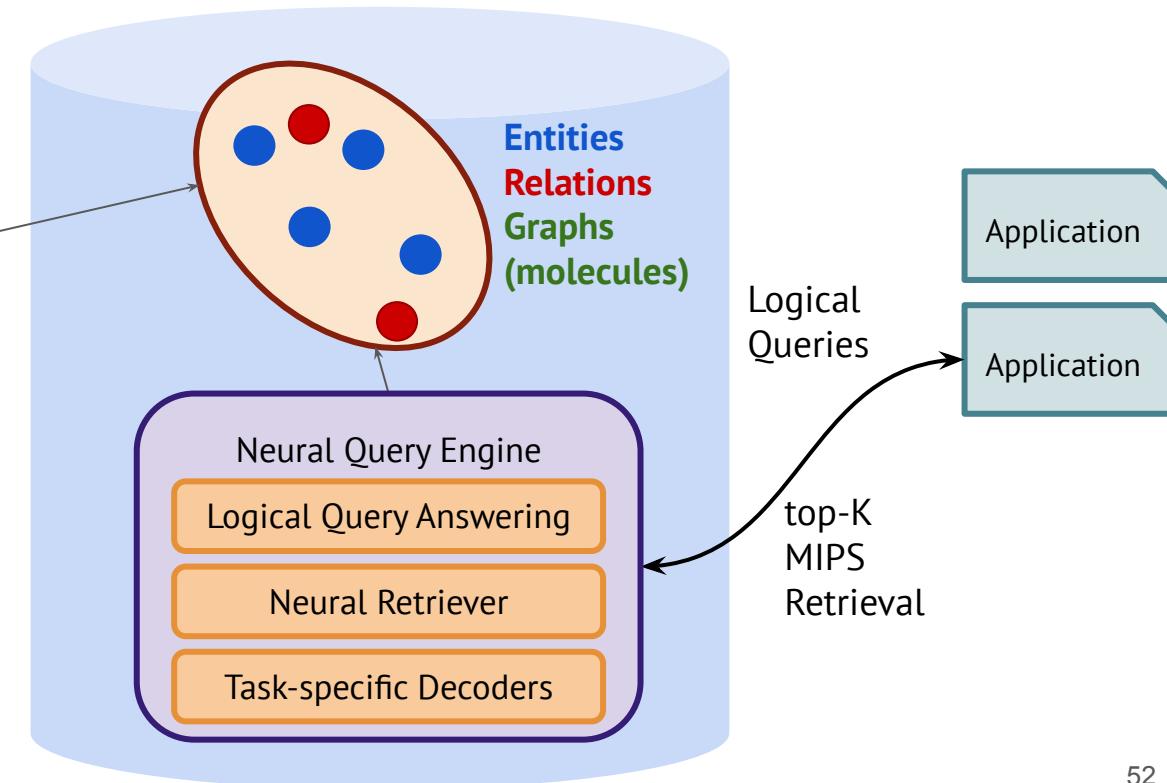
Space of KG Tasks Today



Future: Neural Graph Databases

(Montreal, location, Quebec)
(Quebec, location, Canada)
(Canada, bordersWith, USA)

- No symbolic storage
- Embedding-based storage
- Inferring Missing Links
- Complex Query Answering
- Updatable



Acknowledgements and Papers

1. **Mikhail Galkin**, Priyansh Trivedi, Gaurav Maheshwari, Ricardo Usbeck, Jens Lehmann. **Message Passing for Hyper-Relational Knowledge Graphs.** EMNLP 2020
2. Mehdi Ali, Max Berrendorf, **Mikhail Galkin**, Veronika Thost, Tengfei Ma, Volker Tresp, Jens Lehmann. **Improving Inductive Link Prediction Using Hyper-relational Facts.** ISWC 2021. **Best Research Paper Award**
3. Dimitrios Alivanistos, Max Berrendorf, Michael Cochez, **Mikhail Galkin**. **Query Embedding on Hyper-relational Knowledge Graphs.** ICLR 2022
4. **Mikhail Galkin**, Etienne Denis, Jiapeng Wu, William L Hamilton. **NodePiece: Compositional and Parameter-Efficient Representations of Large Knowledge Graphs.** ICLR 2022



Q&A Time!



@michael_galkin



@mgalkin

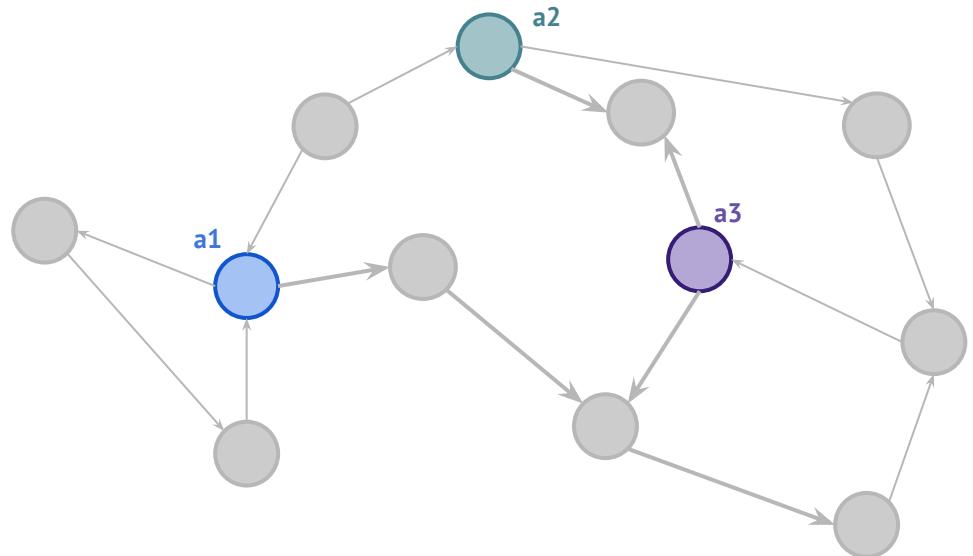


mikhail.galkin@mila.quebec



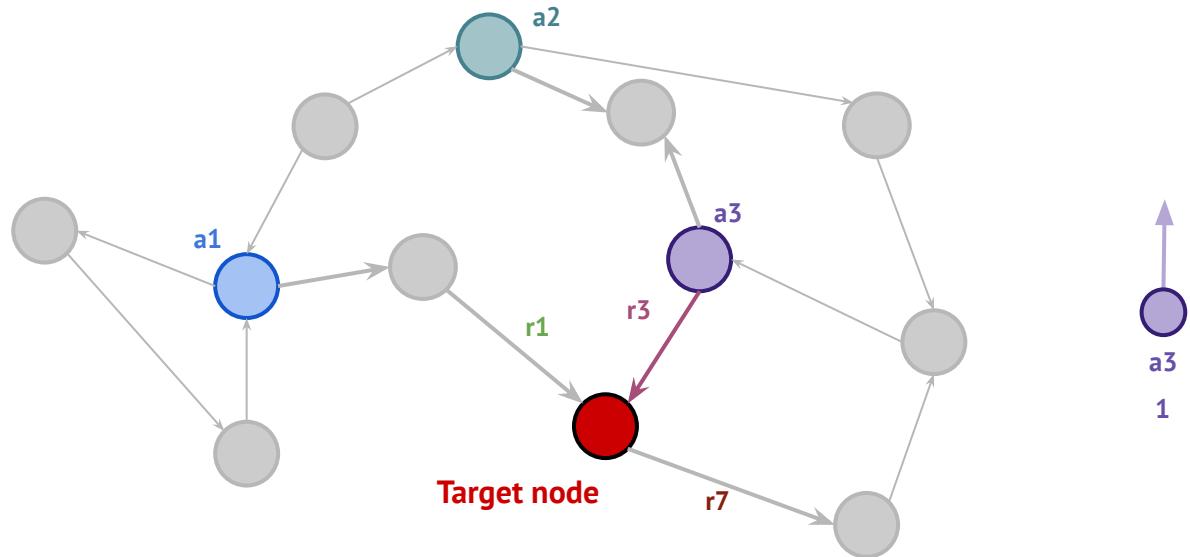
migalkin.github.io

Anchor Node Selection



Current strategy:
40% top degrees
40% top PPR
20% random

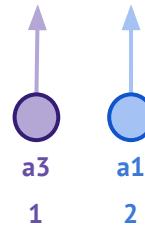
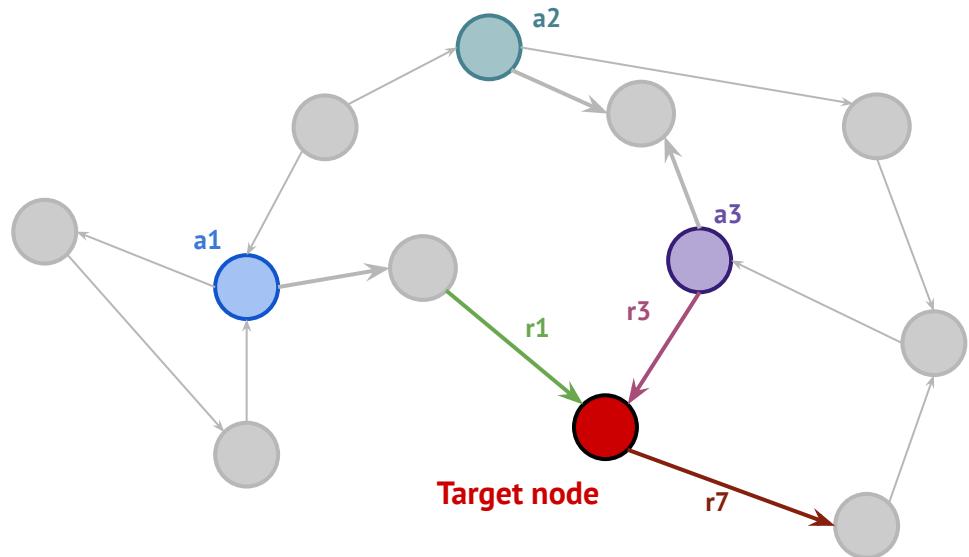
Tokenization



BFS from the target node until we reach $|K|$ anchors

- Can be done in forward pass
- Can be pre-processed and saved

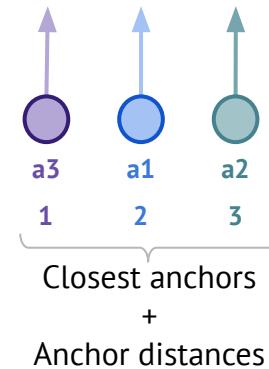
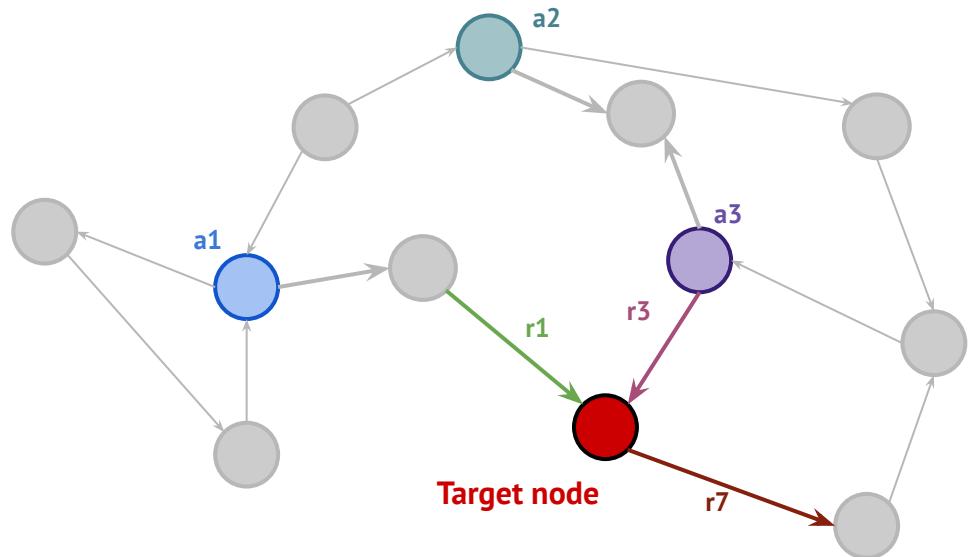
Tokenization



BFS from the target node until we reach $|K|$ anchors

- Can be done in forward pass
- Can be pre-processed and saved

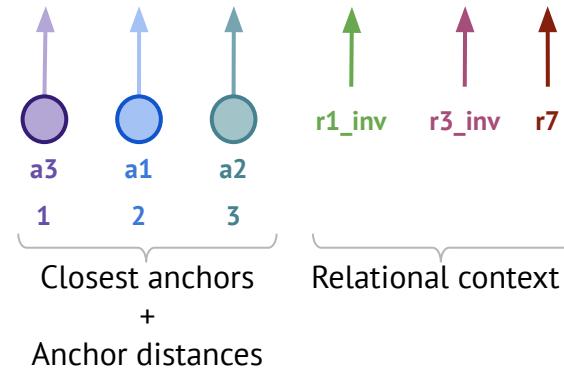
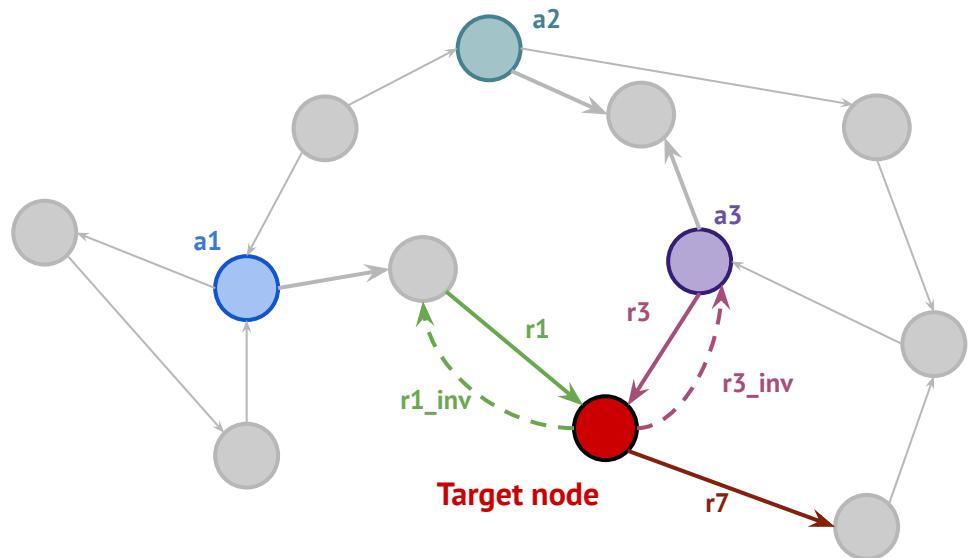
Tokenization



BFS from the target node until we reach $|K|$ anchors

- Can be done in forward pass
- Can be pre-processed and saved

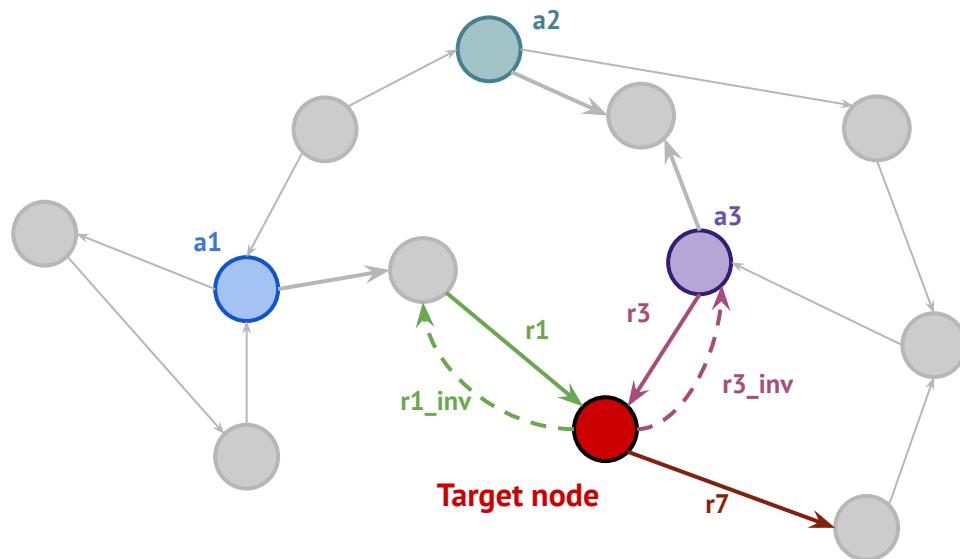
Tokenization



BFS from the target node until we reach $|K|$ anchors

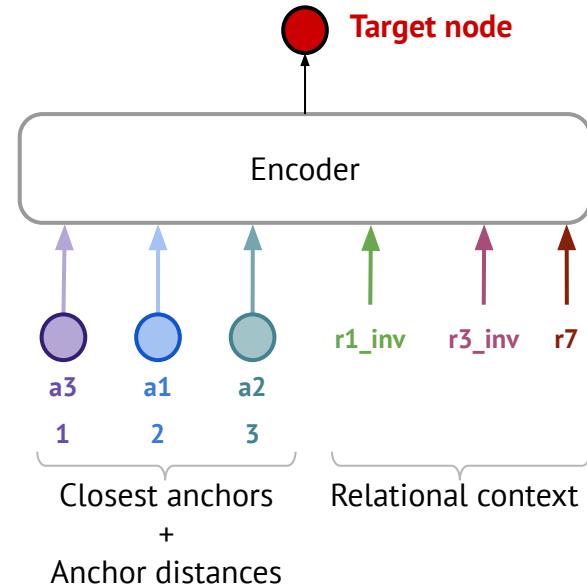
- Can be done in forward pass
- Can be pre-processed and saved

Tokenization

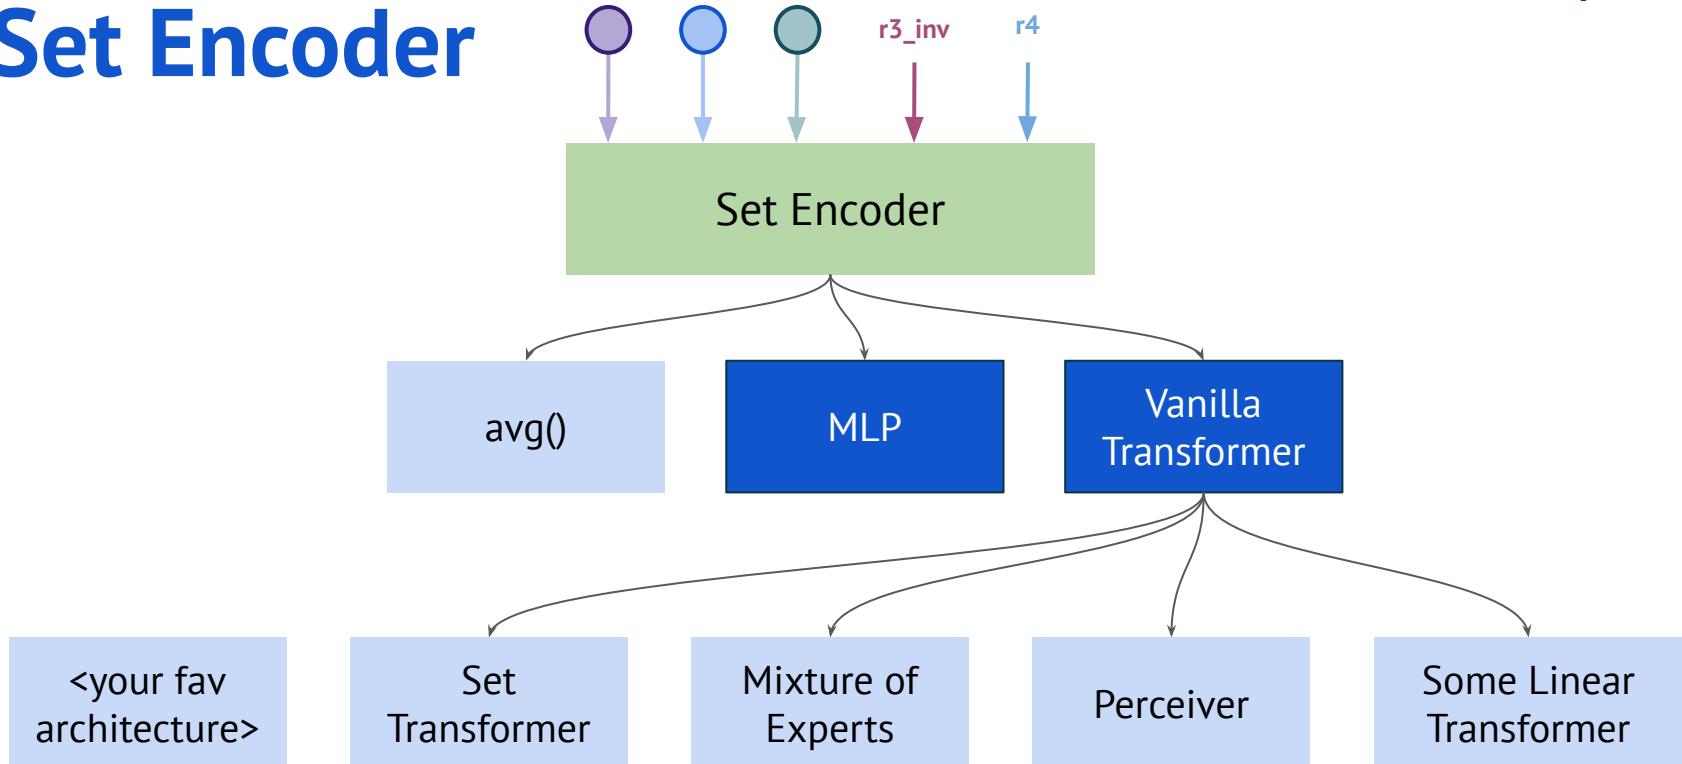


BFS from the target node until we reach $|K|$ anchors

- Can be done in forward pass
- Can be pre-processed and saved



Set Encoder



Transductive Link Prediction

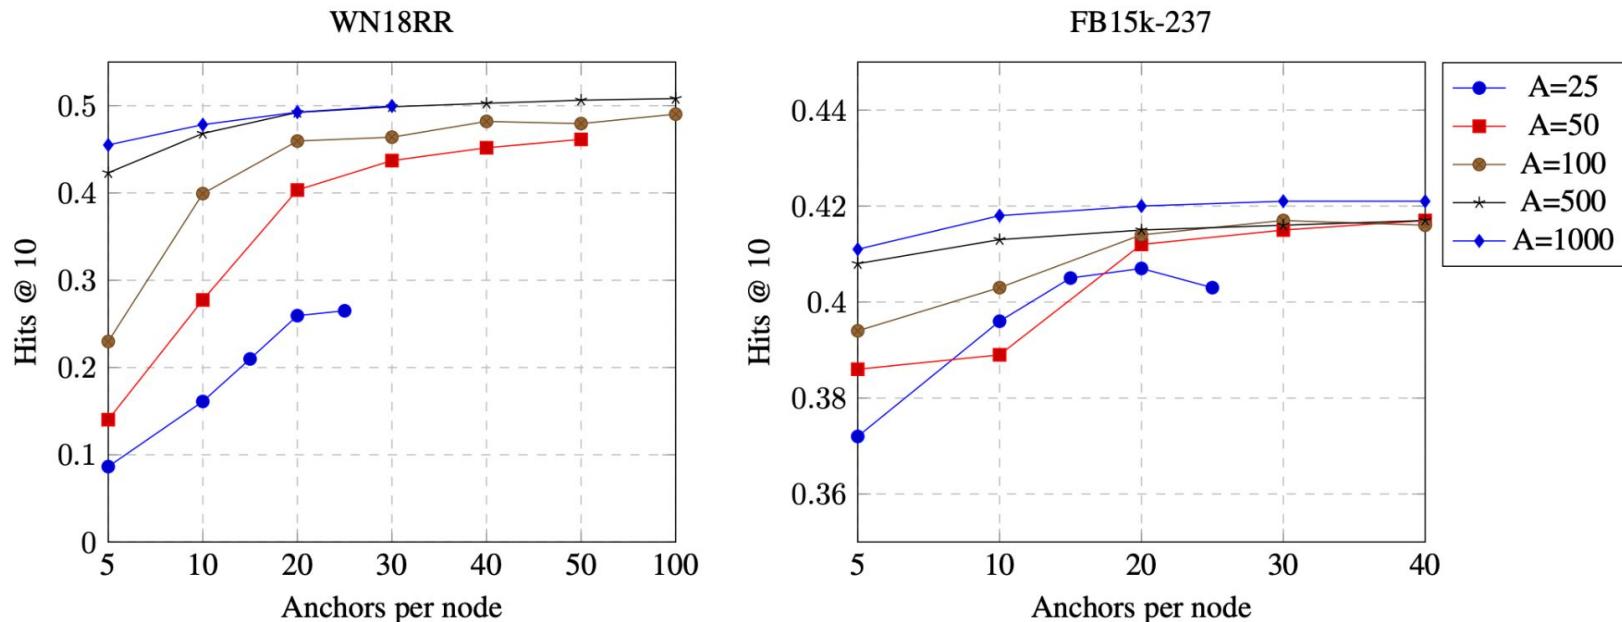


Figure 2: Combinations of total anchors A and anchors per node. Denser FB15k-237 saturates faster on smaller A while sparse WN18RR saturates at around 500 anchors.

NodePiece Experiments: Summary



10x fewer parameters while retaining **90%** of transductive LP



2x better compared to shallow models of similar #params



Relation Prediction and Node Classification: no anchors is better!

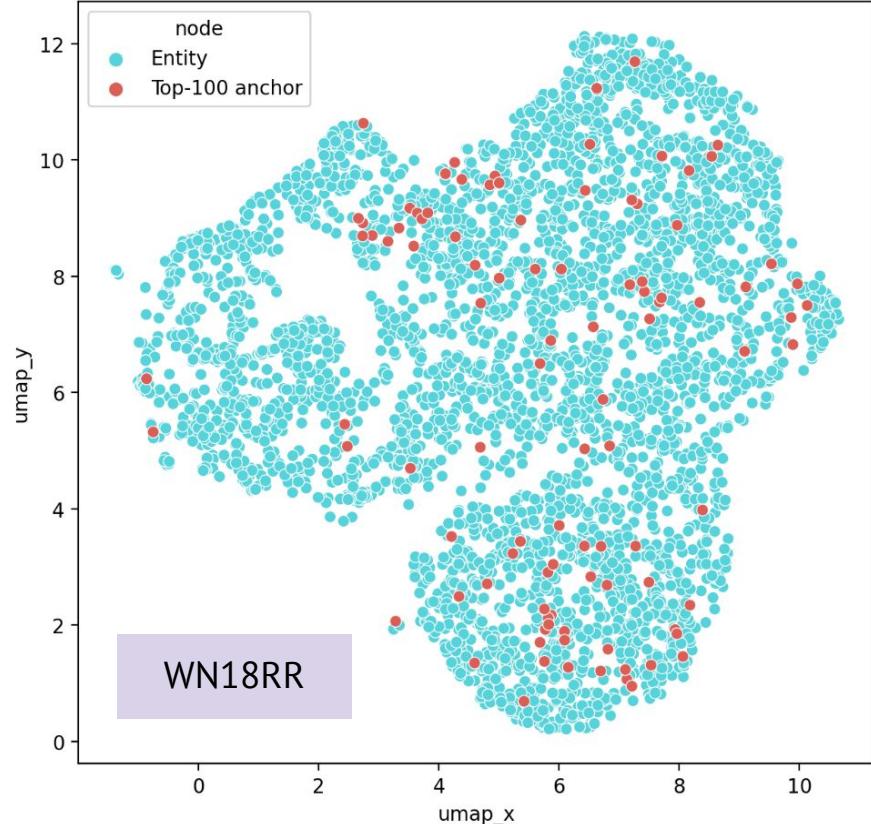
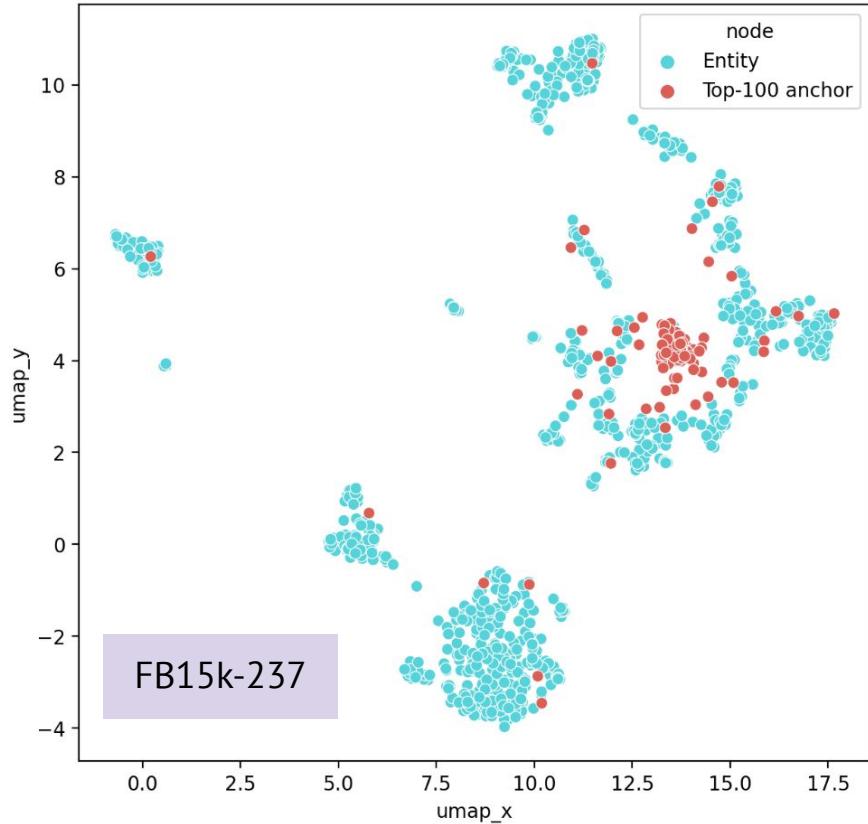


Inductive out-of-the-box and very competitive

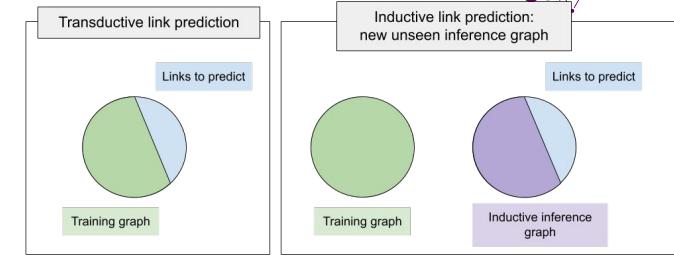
Table 6: Node classification results. $|V|$ denotes vocabulary size (anchors + relations), #P is a total parameter count (millions).

	$ V $	#P (M)	WD50K (5% labeled)			WD50K (10% labeled)		
			ROC-AUC	PRC-AUC	Hard Acc	ROC-AUC	PRC-AUC	Hard Acc
MLP	46k + 1k	4.1	0.503	0.016	0.001	0.510	0.017	0.002
CompGCN	46k + 1k	4.4	0.836	0.280	0.176	0.834	0.265	0.161
NodePiece + GNN	50 + 1k	0.75	0.981	0.443	0.513	0.981	0.450	0.516
- no rel. context	50 + 1k	0.64	0.982	0.446	0.534	0.982	0.449	0.530
- no distances	50 + 1k	0.74	0.981	0.448	0.516	0.981	0.448	0.513
- no anchors, rels only	0 + 1k	0.54	0.984	0.453	0.532	0.984	0.456	0.533

Visualizations: Anchors + Entities



Inductive Link Prediction



Inference graphs are disjoint with training (new nodes)

NodePiece + CompGCN encoder = SOTA on many tasks on relation-rich graphs

Table 5: Inductive link prediction results, Hits@10. Best results are in **bold**, second best are underlined. † results taken from Teru et al. (2020). NBFNet results taken from Zhu et al. (2021).

Class	Method	FB15k-237				WN18RR				NELL-995			
		V1	V2	V3	V4	V1	V2	V3	V4	V1	V2	V3	V4
Path	Neural LP †	0.529	0.589	0.529	0.559	0.744	0.689	0.462	0.671	0.408	0.787	0.827	<u>0.806</u>
	DRUM †	0.529	0.587	0.529	0.559	0.744	0.689	0.462	0.671	0.194	0.786	0.827	<u>0.806</u>
	RuleN †	0.498	0.778	0.877	0.856	0.809	0.782	0.534	0.716	0.535	0.818	0.773	0.614
GNN	GraIL †	0.642	0.818	0.828	0.893	0.825	0.787	0.584	0.734	<u>0.595</u>	0.933	<u>0.914</u>	0.732
	NBFNet	<u>0.834</u>	0.949	0.951	0.960	0.948	0.905	0.893	0.890	-	-	-	-
	NP + CompGCN	0.873	0.939	0.944	0.949	<u>0.830</u>	0.886	<u>0.785</u>	<u>0.807</u>	0.890	<u>0.901</u>	0.936	0.893